
Isogeny Alpha System

Institutional Alpha Framework for Micro E-mini Nasdaq-100 Futures

*An Empirical Performance Study Through Backtesting and Walk-Forward
Validation*



Kairos Capital Research

Cliff Angers

Quantitative Researcher

Research Paper v8.0 • Published June 07, 2026

Instrument: MNQ (Micro E-mini Nasdaq-100 Futures) • Session: 9:30 AM to 12:00 PM ET

Platform: Tradeify \$25,000 Evaluation • Venue: CME Group

Abstract

This paper presents the design, implementation, and empirical performance of the Isogeny Alpha System v8.0, an adaptive, multi-strategy algorithmic trading framework targeting the Micro E-mini Nasdaq-100 (MNQ) futures contract during the U.S. morning trading session (9:30 AM to 12:00 PM ET). The system integrates six complementary intraday strategies Gap Fill, Opening Range Breakout (ORB), Initial Balance (IB) Breakout, VWAP Reversion and Bounce, Fair Value Gap (FVG) fills, and the new 80% Value Area Rule governed by an adaptive regime classifier and a twenty-point institutional confidence scoring layer that dynamically sizes positions and filters low-quality setups before execution.

Version 7.0 delivers three successive upgrade cycles. The Order Flow Upgrade (v6) introduced a two-target exit system (T1 exits 50% at 1R, T2 trails with a 3x intraday ATR Chandelier stop), fixing the critical discovery that 44% of all backtest trades were producing exactly zero P&L; despite averaging 15.7x favorable excursion. Four new order flow signals were added: time-of-day adjusted RVOL (hard block below 0.8x), Wyckoff absorption detection, Kyle's lambda informed flow proxy, and CVD climax/exhaustion detection. The Research Upgrade (v7) further expanded the scoring to 20 points by adding SMH semiconductor lead signal, CFTC COT Leveraged Funds positioning, Anchored VWAP proximity, and daily market breadth. The HMM was upgraded from 3 to 5 states with a multivariate feature set. A walk-forward validation framework confirmed the system's robustness with a Walk-Forward Efficiency of 201%, indicating out-of-sample performance exceeds in-sample performance.

The v7.1 Hybrid System backtested across 623 trading days (June 2024 to June 2026) on Databento GLBX.MDP3 NQ continuous futures produced 283 trades with a 67.1% win rate and net P&L; of \$6,794 — averaging \$2,265 per year. Maximum drawdown was \$438.82, well within the \$1,000 Tradeify limit. The system was profitable in all 3 calendar years tested: 2024 (65.4% WR, +\$2,036), 2025 (66.0% WR, +\$2,942), and 2026 partial (73.1% WR, +\$1,816). The out-of-sample walk-forward validation produced a 71.4% win rate and \$808 P&L; on data the system had never seen, confirming the edge is structural, not overfit.

Version 8.0 builds directly on the v7.1 hybrid engine and adds an optional LightGBM meta-labeling layer (Section 14.5): a secondary classifier that observes the rule-based engine's own candidate trades and learns to veto the weakest of them, never to create or override a signal. Four of five trained per-strategy models cleared a two-bar deployability gate (purged-CV out-of-sample AUC above 0.55 and expanding walk-forward efficiency at or above 50%) and are certified live-eligible; the fifth was automatically excluded after its edge was shown to be curve-fit. Run as a single combined engine exactly as it would trade live, the ML-gated system produced 220 trades at an 80.9% win rate and \$8,177 net P&L; over the identical two-year window, a 23% reduction in trade

count alongside a 25% increase in profit and a 15% reduction in maximum drawdown versus the rules-only baseline.

Win Rate	Net P&L;	Max Drawdown	Avg R:R	Avg/Year
67.1%	\$6,794	\$438	4.77x	\$2,265
190W / 93L of 283	2-year Databento BT	vs \$1,000 limit	two-target exit	3/3 positive years

Table of Contents

1	Introduction & Research Motivation	4
2	Market Context: NQ Futures & Prop Firm Evaluation	5
3	System Architecture Overview	7
4	Market Regime Classification Framework	8
4.1	EMA Trend Detection	8
4.2	Adaptive ATR Volatility	9
4.3	VIX Regime Gating	10
4.4	Direction Allowance Logic	10
4.5	Overnight Range & Day Type Classification	11
4.6	VIX Term Structure & VVIX Vol-of-Vol Gate	11
4.7	Open Type Classifier (CME Auction Theory)	12
4.8	0DTE / Expiry Day Gamma Context	12
5	Strategy Modules	13
5.1	Gap Fill	13
5.2	Opening Range Breakout (ORB)	15
5.3	Initial Balance Breakout (IB)	18
5.4	VWAP Reversion & Bounce	20
5.5	Fair Value Gap (FVG)	23
6	Trade Simulation Engine	25
7	Risk Management Framework	27
8	Backtesting Methodology	30
9	Performance Results	34
9.1	Overall Statistics Three-Way Comparison	34
9.2	Per-Strategy Breakdown	35
9.3	Confidence Score Distribution	36
9.4	Regime Analysis	37
9.5	Day-of-Week Analysis	38
9.6	Drawdown Analysis	39
10	Live Implementation	41
10.1	Real-Time Monitor Architecture	41

10.2	Direction Lock & Contradictory Signal Prevention	42
10.3	Trade Confirmation Flow (y/n)	42
10.4	Signal Pipeline Latency	43
10.5	Notification System	43
11	Institutional Signal Overlay 20-Point Scoring System	44
11.1	Order Flow Imbalance (OFI) + CVD Divergence	44
11.2	VPIN Toxicity Gate	45
11.3	GEX Gamma Exposure Regime	46
11.4	Hidden Markov Model 5-State Upgrade	46
11.5	Time-Series Momentum & Session Conviction	47
11.6	XLK/SPY Sector Relative Strength + SMH Lead Signal	47
11.7	DXY + TNX Macro Headwind/Tailwind + COT Positioning	48
11.8	NQ/ES Spread Divergence	48
11.9	PDH/PDL/PMH/PML Key Levels	49
11.10	Volume Profile POC, VAH, VAL, Naked VPOC, Composite	49
11.11	HAR-RV Stop Multiplier	50
11.12	RVOL Time-of-Day Adjusted Relative Volume	51
11.13	Absorption Detection (Wyckoff Effort vs Result)	51
11.14	CVD Climax / Exhaustion Signal	52
11.15	Opening Candle Continuation (OCC)	52
11.16	Kyle's Lambda Informed Flow Proxy	53
11.17	Anchored VWAP Yearly, Swing Low, Weekly	53
11.18	Market Breadth QQQ/IWM RS + \$ADDN	53
11.19	Complete 20-Point Confidence Scoring System	54
12	Regime-Contextual Bot Memory & Adaptive Scoring	56
12.1	Signal Logging Before User Confirmation	56
12.2	Trade Confirmation and Outcome Tracking	56
12.3	Regime-Contextual Win Rate Learning	57
12.4	Adaptive Confidence Score Adjustment	57
13	Order Flow Upgrade Two-Target Exit System	58
13.1	The Breakeven Problem: 44% of Trades Were \$0 P&L;	58
13.2	T1/T2 Two-Target Architecture	59

13.3	Chandelier Trailing Stop (T2)	59
13.4	Strategy-Specific Target Extensions	60
13.5	80% Value Area Rule New Strategy	60
13.6	Single Print Zones as Structural Targets	61
14	Walk-Forward Validation	62
14.1	Methodology and WFE Ratio	62
14.2	Results and Robustness Assessment	62
14.5	Machine Learning — Meta-Labeling Layer	63
14.5.1	Architecture: Rules Primary, ML Veto-Only	63
14.5.2	The Two-Bar Deployability Gate	63
14.5.3	Live Results: Two-Year Combined Backtest	64
14.5.4	Walk-Forward Confirmation	64
14.5.5	Interpretability: SHAP & Decision Surfaces	65
14.5.6	Operational Recommendation	65
15	TradingView Pine Script Integration	66
16	Limitations & Risk Factors	65
17	Conclusion	67
Appendix A	Strategy Parameter Reference	69
Appendix B	Regime Gate Summary	70
Appendix C	Prop Firm Compliance Checklist	71
Appendix D	Institutional Module Parameter Reference	72
Appendix E	Glossary of Terms	73

How To Read This Paper

This paper was written for two types of readers simultaneously: the experienced quantitative analyst who wants the mathematical rigor, and the complete beginner who just wants to understand what is actually happening and why it works. Both readers will find everything they need here you do not need to skip anything.

If you are new to trading or quantitative finance:

Every technical term is defined when it first appears. Every mathematical formula is followed immediately by a plain-English explanation of what it is actually saying. Every strategy is explained with a real example using real NQ prices. Look for these visual boxes throughout the paper:

BLUE BOX	Plain English: A simple explanation of the concept just described above it.
GREEN BOX	Real Example: An actual trade walkthrough with real NQ prices and dollar amounts.
ORANGE BOX	Important: A critical warning or key insight you absolutely must not miss.
	Key Term: A precise definition of a trading or mathematics term.

If you are an experienced quant or trader:

The mathematical sections, academic citations, formula derivations, and statistical analysis are complete and rigorous. You can skim the plain-English boxes and move directly through the formulas, tables, and results. All parameter choices are justified with published research. The backtest methodology follows standard walk-forward validation with an embargo window.

Important: Before You Start

This paper describes a real trading system that was built and tested on a real \$25,000 prop firm evaluation account. The numbers are real. The trades are real (in backtest). The math is real. But trading always involves risk of loss. Nothing in this paper is financial advice. The system described here is a decision-support tool not a guarantee of profit.

1. Introduction & Research Motivation

Systematic intraday trading of equity index futures presents a well-documented opportunity for consistent edge extraction when grounded in empirically validated market microstructure principles. The Nasdaq-100 futures complex, specifically the Micro E-mini (MNQ, \$2/point) offers retail-accessible leverage, deep institutional participation, and well-defined intraday session structure that creates repeatable, quantifiable patterns.

Plain English: What does this opening paragraph mean?

Fancy way of saying: there are real, documented patterns in the Nasdaq futures market that repeat themselves enough that a computer program can find them, trade them, and make money. This paper explains exactly what those patterns are, why they work, and how the system identifies them automatically so you just have to decide whether to press the button.

1.1 What Is a Futures Contract? (Start Here If You Are New)

A futures contract is an agreement to buy or sell something at a specific price on a specific future date. In this case, the 'something' is the Nasdaq-100 stock index a list of the 100 largest tech and growth companies listed on the Nasdaq exchange (Apple, Microsoft, Nvidia, Amazon, Google, Meta, Tesla, and others).

You are not buying actual shares of those companies. You are making a bet on whether the combined value of those 100 companies goes up or down over the next few hours. If you think they will go up, you 'go long' (buy). If you think they will go down, you 'go short' (sell). At the end of the trade, you settle in cash no actual shares change hands.

KEY TERM | Futures Contract

A financial agreement to buy or sell an asset at a predetermined price on a predetermined date. For equity index futures (like NQ), you are trading on the price movement of a stock market index, not buying the actual stocks. You profit if the index moves in the direction you predicted, and lose if it moves against you.

1.2 What Is the NQ Contract Specifically?

NQ is the ticker symbol for the E-mini Nasdaq-100 futures contract traded on the Chicago Mercantile Exchange (CME). 'E-mini' simply means it is an electronically traded, smaller version of the original floor-traded contract. Each NQ contract represents \$20 per index point. The system in this paper trades MNQ (Micro E-mini), which is \$2 per index point exactly one-tenth the size of NQ, making it accessible for smaller account sizes.

Real Example: NQ vs MNQ Contract Sizes

NQ (E-mini): \$20 per index point standard contract, ~\$400,000 notional value

MNQ (Micro): \$2 per index point 1/10th size, ~\$40,000 notional value

Example: NQ moves from 20,000 to 20,050 (up 50 points)

NQ profit: 50 points x \$20 = \$1,000

MNQ profit: 50 points x \$2 = \$100

The system uses MNQ because \$2/point means less dollar risk per trade, which is critical for prop firm evaluations with strict drawdown limits.

1.3 What Is a Prop Firm Evaluation?

A proprietary trading firm (prop firm) gives traders access to a large funded trading account in this case \$25,000 in exchange for a cut of the profits. But first, the trader must pass an 'evaluation' to prove they can trade profitably without blowing up the account. The Tradeify \$25,000 evaluation has three rules:

Rule	Limit	What It Means in Plain English
Profit Target	\$1,500	You must make at least \$1,500 net profit. There is no time limit take as many days as you need.
Trailing Drawdown	\$1,000 max	Your account can never fall more than \$1,000 below its highest end-of-day balance. This is the hardest rule once you make money, that money becomes part of the floor.
Consistency Rule	No day > 40% of total profit	No single day can account for more than 40% of all your profits. You cannot make \$1,500 in one lucky day and call it done.

Plain English: The Trailing Drawdown The Most Important Rule

Here is why the trailing drawdown rule is so dangerous if you do not respect it. You start with \$25,000. You make \$800 profit great, you now have \$25,800. But NOW your floor rises to \$24,800 (always \$1,000 below your highest balance). If you then lose \$900, your balance is \$24,900 which is still above \$24,800, so you survive. But if you lose \$1,100 from that \$25,800 high, your balance falls to \$24,700, which is below the \$24,800 floor evaluation FAILED. The more you make early, the higher the floor rises, and the less room you have to lose. This is why the system is designed to be LOW-VARIANCE first and profitable second.

1.4 What Is an Algorithmic Trading System?

An algorithmic (or 'algo') trading system is a computer program that monitors the market continuously, identifies specific patterns or conditions that historically precede profitable price moves, and alerts the trader (or executes automatically) when those conditions are met. The system removes emotion from trading it does not panic, does not get greedy, and does not change its mind based on news headlines.

The Isogeny Alpha System is a semi-automated system: the computer does all the analysis and generates all the signals, but the human operator makes the final decision to enter each trade. This keeps the trader in control of risk while eliminating the emotional decision-making that destroys most manual traders.

KEY TERM | Algorithm

A precise set of rules that a computer follows to solve a problem or make a decision. In trading, an algorithm defines exactly what market conditions must be present before a trade signal is generated for example, 'price must close above the opening range high AND volume must be 1.5x normal AND VIX must be below 25 AND the trend must be bullish.'

1.5 Why Quantitative Methods?

The word 'quantitative' simply means we use numbers and mathematics instead of subjective judgment. A quantitative trader asks: 'In the past 2,000 trading sessions, when X happened, what was the probability that Y happened next?' They measure everything, test everything, and only trade patterns that have statistically proven themselves across large historical datasets not patterns that 'look good' on a chart.

This research was motivated by the requirements of the Tradeify \$25,000 evaluation program. The constraints are intentionally tight, rewarding low-variance, high-win-rate approaches over high-volatility speculative strategies. Quantitative methods are ideal for this because they allow us to precisely measure and control both the win rate and the variance before risking any real money.

1.6 Design Principles

The Isogeny Alpha System was designed around five core principles:

- **Empirical grounding:** every strategy is anchored to published research with documented win rates on ES/NQ futures across multi-year datasets. No strategy was included because it 'looks good' every one has a mathematical foundation.
- **Adaptive regime awareness:** static parameters are replaced with ATR-normalized dynamic thresholds that self-adjust to current volatility. The same system works in a VIX 12 calm grind and a VIX 40 crisis.
- **Defense-first risk model:** the \$50 maximum risk per MNQ trade (25 points x \$2) means even a worst-case day of 3 full losses only costs \$150 well within the \$1,000 trailing drawdown limit.
- **Minimal discretion:** all signal generation, filtering, and risk checks are algorithmic. Human judgment is limited to the binary decision of whether to take a generated signal.
- **Live-ready implementation:** the system runs a real-time bar cache with signal latency under 500ms, macOS push notifications, and a trade journal not just a backtest.

1.7 Scope of this Paper

This document covers the complete system from first principles: market context, regime classification, six strategy modules, the two-target exit system, trade simulation methodology, all 20 institutional confidence signals, backtesting results, walk-forward validation, and live implementation architecture. Every concept is explained from scratch so that a reader with no prior trading or mathematics background can understand and evaluate the system independently.

2. Market Context: NQ Futures & Prop Firm Evaluation

Before diving into the system itself, we need to understand the instrument being traded and the environment in which it is traded. This section explains futures contracts, the MNQ specification, how leverage works, and why the prop firm evaluation model is uniquely suited to a systematic approach like this one.

2.1 The MNQ Contract

The Micro E-mini Nasdaq-100 futures contract (ticker: MNQ) was introduced by CME Group in May 2019 as a 1/10th-sized version of the standard NQ contract. Key specifications:

Specification	Value
Contract Multiplier	\$2.00 per point
Tick Size	0.25 index points (\$0.50 per tick)
Typical Daily Range	120 to 250 points (\$240 to \$500 per contract)
Trading Hours	Sunday 6 PM to Friday 5 PM ET (23 hours/day)
Primary Session	9:30 AM to 4:00 PM ET
Margin (Tradeify, 1 micro)	~\$50 intraday
Average Daily Volume	~450,000 contracts
Correlation with QQQ	>0.99

Plain English: Reading the Contract Spec Table

CONTRACT MULTIPLIER = \$2/point means: if NQ moves up 10 points and you are long (bought), you make $10 \times \$2 = \20 . If NQ moves down 10 points and you are long, you lose \$20. TICK SIZE = 0.25 points means the smallest price movement you will ever see is 0.25 points, worth \$0.50. So NQ goes: 20,000.00... 20,000.25... 20,000.50... 20,000.75... 20,001.00 It cannot jump from 20,000 to 20,000.30 it always moves in 0.25 increments.

2.1.1 How Leverage Works (Critically Important)

Futures trading involves leverage you control a large position with a small amount of actual money. At NQ = 20,000 points, one MNQ contract has a 'notional value' of $20,000 \times \$2 = \$40,000$. But you only need a few hundred dollars in margin to hold it. This is both the power and the danger of futures.

Real Example: Leverage in Action Both Directions

NQ is at 20,000. You buy (go long) 1 MNQ contract.

Notional value: $20,000 \times \$2 = \$40,000$

Margin required: ~\$200-500 (varies by broker)

NQ rises 50 points to 20,050:

Your profit: 50 points x \$2 = \$100 (on a \$200-500 margin that's a 20-50% return!)

NQ falls 25 points to 19,975:

Your loss: 25 points x \$2 = \$50 (the system's max stop = 25 points = \$50 max loss)

KEY INSIGHT: The system ALWAYS uses a stop loss of maximum 25 points = \$50 max loss.

No matter what NQ does, one trade cannot lose more than \$50.

Important: Leverage is a Double-Edged Sword

The same leverage that makes futures attractive also makes them dangerous if you do not manage risk. A 25-point move against you on 1 MNQ costs \$50. But on 10 MNQ contracts, the same 25-point move costs \$500. The system NEVER trades more than 2 MNQ contracts at once, keeping maximum trade risk at \$100. This is not a limitation it is the entire reason the system can survive a losing streak without failing the evaluation.

At current NQ levels (~20,000), each full point of movement equals \$2. The system's maximum stop of 25 points represents a maximum loss of \$50 per trade, meaning even a catastrophic streak of 20 consecutive losses would only produce a \$1,000 drawdown, exactly at the Tradeify limit. In practice, the 76.7% win rate means the probability of 10+ consecutive losses is less than 0.001%.

2.2 Tradeify Evaluation Rules

The Tradeify \$25,000 evaluation operates under a trailing drawdown model, unlike traditional fixed drawdown accounts, the floor rises as the end-of-day (EOD) balance grows. This creates a unique constraint: early profits both help (higher balance) and hurt (higher floor to protect). Key rules:

Rule	Limit	System Response
Profit Target	\$1,500 net profit	Target met in backtest by Day 42 average
Trailing Drawdown	\$1,000 from highest EOD balance	Max simulated DD: \$300 (30% utilization)
Daily Loss Limit	No explicit limit (Tradeify)	\$100 self-imposed hard stop
Max Contracts	1 mini / 10 micros	System trades 1 MNQ (expandable to 2 at high WR)
Consistency Rule	No day > 40% of total profit	Max 2 trades/session limits daily upside
Trading Days	Minimum 5 days required	Average 3.2 trades per 5-day week

News Events	Prohibited on FOMC/NFP days	VIX gate >25 effectively blocks these days
-------------	-----------------------------	--

Key insight: The trailing drawdown model means aggressive early trading is MORE dangerous than cautious trading. A trader who wins \$800 on Day 1 now has a floor of \$24,800, losing it back over Days 2 to 3 fails the evaluation. The system's low-variance approach (avg 1.2 trades/day, \$50 max risk) is specifically designed for this structure.

3. System Architecture Overview

The Isogeny Alpha System consists of five interconnected layers, each with a distinct responsibility:

Layer	Module	Function
Data	data_loader.py / yfinance	5-min OHLCV for NQ=F; VIX, VIX3M, VVIX, XLK, SPY, DXY, TNX daily closes
Regime	quant_regime.py	EMA trend, adaptive ATR, VIX regime, overnight range type, VIX term structure, VVIX gate, open-type classifier, expiry context
Signal	quant_gap / orb / ib / vwap / fvg .py	Five independent strategy detectors; each returns a typed signal or None
Inst. Filters	inst_ofi / vpin / gex / hmm / harv / tsmom / leadlag / levels / sectors / macro / volprofile .py	12 institutional signals: hard blocks (BNS, OFI, VVIX, backwardation, macro) + soft scoring (CVD, overnight, term structure, sector, NQ/ES spread, conviction)
Hybrid Engine	hybrid_engine.py	12-point confidence scorer + HAR stop multiplier; skips score <= 3; sizes 2 contracts at score >= 10
Monitor	monitor.py / fast_feed.py	Direction lock (no contradictory signals); y/n trade confirmation; PDH/PDL/PMH/PML level alerts; bot memory integration
Memory	bot_memory.py	Logs every real signal before confirmation; tracks outcomes; regime-contextual WR learning; adjusts per-strategy confidence score
Notifications	notifications.py	macOS popup + sound within 200ms of signal detection

Plain English: How the Layers Work Together (Simple Version)

Think of the system as a factory assembly line with quality control at every stage. Raw material (price data) enters at the left. Each station adds value and passes only the best product to the next station. By the end, only the highest-quality trade signals make it all the way through to the notification that reaches your phone. Most signals get filtered out somewhere along the line and that is exactly the point.

THE SIGNAL PIPELINE FROM DATA TO NOTIFICATION

- 1

Data Layer Downloads 5-minute NQ bars, VIX/VIX3M/VVIX, XLK/SPY/SMH, DXY/TNX, COT data. All pre-loaded at 9:20 AM startup. Zero network calls during trading.
- 2

Regime Layer Classifies the current market: trend direction (strong_bull/neutral/bear), volatility regime (VIX level), HMM state (5-state latent regime), overnight range, VIX term structure.
- 3

Strategy Layer Six detectors run in priority order: Gap Fill, FVG, ORB, IB Breakout, VWAP Rev/Bounce, VA Rule. Each either returns a signal or None.
- 4

Hard Block Layer 11 filters that IMMEDIATELY reject a signal regardless of score: BNS jump, OFI opposing, RVOL thin, CVD climax, absorption wall, VPIN high, large gap, macro headwind, VIX crisis.

5	20-Point Scorer If not blocked, score the signal across 20 institutional dimensions. Skip if score ≤ 5 . Use 1 MNQ if 6-15. Use 2 MNQ if ≥ 16 .
6	Notification Fire macOS popup and sound. Wait for y/n confirmation. Only confirmed trades count toward the 3-trade daily limit.

The core design principle is **separation of concerns**: the regime layer knows nothing about individual strategies; each strategy knows nothing about the others; the engine composes them in priority order. This allows any single module to be improved, replaced, or removed without breaking the system.

3.1 Strategy Priority Order

When multiple strategies are valid on the same day, the engine applies them in priority order until the daily trade limit (3) is reached or the daily loss limit (\$100) is triggered:

Priority	Strategy	Rationale
1	Gap Fill	Highest WR (78%+), fires at open, captures institutional flow
2	FVG	Works in any regime; high WR on quality setups; fires from 9:45
3	ORB	Strong trend continuation; pullback entry improves R:R significantly
4	IB Breakout	Fires 10:00 to 11:30 only; requires confirmed IB range + C-period
5	VWAP Rev	Lower WR; only fires in neutral trend + normal volatility
5	VWAP Bounce	Trend continuation at VWAP; AM and PM windows

4. Market Regime Classification Framework

The regime layer is the engine's nervous system. Every trading decision is contingent on the current market state. Unlike static systems that use fixed lookback periods, the NQ Quant System uses three orthogonal dimensions: EMA trend strength, ATR-normalized volatility, and VIX absolute level.

4.1 EMA Trend Detection

The first thing the system does every morning is classify the current market trend. It does this by comparing two Exponential Moving Averages (EMAs) of NQ's daily closing price: a fast 8-day EMA and a slow 21-day EMA. The relationship between these two lines tells the system which direction the market has been moving recently.

KEY TERM | Exponential Moving Average (EMA)

A running average of prices that gives more weight to recent prices than older prices. Unlike a Simple Moving Average (SMA) which weights all days equally, the EMA responds faster to recent price changes. The '8-day EMA' gives the most weight to the last 8 days of prices. The '21-day EMA' gives the most weight to the last 21 days of prices. When the fast EMA (8-day) is above the slow EMA (21-day), the recent trend is up (bullish). When it is below, the recent trend is down (bearish).

Trend direction is measured by the spread between an 8-day and 21-day exponential moving average of daily closing prices, expressed as a percentage of the slow EMA. The EMA is computed recursively with smoothing factor alpha:

$$EMA_n(t) = \alpha \cdot P(t) + (1 - \alpha) \cdot EMA_n(t - 1), \text{ where } \alpha = 2 / (n + 1) \quad (1)$$

In plain English: Today's EMA = (a small fraction of today's price) + (a large fraction of yesterday's EMA). The 'alpha' controls how fast it responds: for the 8-day EMA, $\alpha = 2/9 = 0.222$, meaning today's price gets 22.2% of the weight and yesterday's EMA gets 77.8%. For the 21-day EMA, $\alpha = 2/22 = 0.091$, so it responds much more slowly to price changes.

$$\text{strength} = [EMA_8(t) - EMA_{21}(t)] / EMA_{21}(t) \times 100 \quad (2)$$

In plain English: Subtract the slow average from the fast average, divide by the slow average, multiply by 100 to get a percentage. If the 8-day EMA is 20,200 and the 21-day EMA is 20,000, the strength is $(20,200 - 20,000) / 20,000 \times 100 = +1.0\%$, meaning the recent trend is mildly bullish. If the 8-day EMA is 19,600 and the 21-day EMA is 20,000, strength = -2.0% , meaning the recent trend is moderately bearish.

Strength Range	Classification	Long Allowed	Short Allowed	Frequency*
strength > +3%	STRONG BULL	Yes Preferred	No Blocked	~18%
+1% < strength <= +3%	bull	Yes Yes	Limited	~22%
-1% <= strength <= +1%	neutral	Yes Yes	Yes Yes	~30%

-3% <= strength < -1%	bear	Limited	Yes Yes	~20%
strength < -3%	STRONG BEAR	No Blocked	Yes Preferred	~10%

**Approximate frequency during 2020-2025 NQ daily data*

Real Example: Real EMA Classification Examples

April 2025 (tariff shock): 8-day EMA = 18,200, 21-day EMA = 19,500

strength = $(18,200 - 19,500) / 19,500 \times 100 = -6.7\%$ --> STRONG BEAR

System: only trades short setups, blocks all long entries

May 2025 (recovery rally): 8-day EMA = 20,100, 21-day EMA = 19,700

strength = $(20,100 - 19,700) / 19,700 \times 100 = +2.0\%$ --> BULL

System: prefers long setups, allows carefully filtered short setups

Normal sideways market: 8-day EMA = 19,850, 21-day EMA = 19,810

strength = $(19,850 - 19,810) / 19,810 \times 100 = +0.20\%$ --> NEUTRAL

System: all 6 strategies active, trades both long and short

The EMA8/EMA21 combination was chosen for its responsiveness to regime changes without excessive noise. The 3% strong threshold captures only genuine sustained trends in a VIX 30+ crash environment, EMA spread can reach 8-12%, making the strong classifications clearly visible and valid. The 1% mild threshold prevents misclassification during normal daily fluctuations.

4.2 Adaptive ATR Volatility How the System Measures Market Noise

Every trading session is different. Some days NQ moves 80 points total. Other days it moves 400 points. If you set a stop loss of 20 points on a 400-point day, it will be hit by random noise before the trade has a chance to work. If you set a stop of 20 points on an 80-point day, you might be risking more than the whole day's range. The system solves this with the Average True Range (ATR) a dynamic measure of how much the market is actually moving.

KEY TERM | Average True Range (ATR)

A measure of market volatility that tells you, on average, how much the market moves per day (or per bar). 'True Range' for each day is the largest of three values: (1) today's High minus today's Low, (2) the absolute difference between today's High and yesterday's Close, (3) the absolute difference between today's Low and yesterday's Close. ATR averages these True Range values over N days. A higher ATR means a more volatile market.

Rather than using a fixed ATR period, the system takes the maximum of the 5-day and 20-day ATR computed from daily high-low ranges using Wilder's smoothing method:

$$ATR_n(t) = [(n - 1) \cdot ATR_n(t - 1) + TR(t)] / n \quad (3)$$

In plain English: Today's ATR = (yesterday's ATR times a big fraction) plus (today's True Range times a small fraction). For a 5-day ATR: $ATR = (4/5 \times \text{yesterday's ATR}) + (1/5 \times \text{today's range})$. This gives a smoothly evolving estimate of volatility that does not jump wildly day-to-day.

$$ATR_{\text{adaptive}} = \max(ATR_5, ATR_{20}) \quad (4)$$

In plain English: Take WHICHEVER is larger: the 5-day ATR or the 20-day ATR. During a volatility spike, the 5-day ATR spikes first and makes the system automatically widen stops to avoid being stopped out by noise. During a calm recovery, the 20-day ATR remains elevated (it includes the spike in its average) and prevents the system from getting too aggressive too soon.

Real Example: ATR in Practice Calm vs Volatile Market

CALM MARKET (April 2024, VIX ~14):

5-day ATR = 120 pts 20-day ATR = 130 pts

Adaptive ATR = $\max(120, 130) = 130$ pts

ORB min range = $130 \times 0.025 = 3.25$ pts (tiny moves filtered out)

VWAP stop = $130 \times 0.06 = 7.8$ pts (tight, appropriate for calm day)

VOLATILE MARKET (April 2025, tariff shock, VIX ~35):

5-day ATR = 380 pts 20-day ATR = 210 pts

Adaptive ATR = $\max(380, 210) = 380$ pts

ORB min range = $380 \times 0.025 = 9.5$ pts (filters out small fake breakouts)

VWAP stop = $380 \times 0.06 = 22.8$ pts (wider stop needed in chaos)

The system AUTOMATICALLY adjusts its parameters to current conditions.

You never need to manually change stop distances based on market conditions.

The max operator ensures: (1) during a volatility spike, the 5-day ATR captures the spike and widens stops/filters appropriately; (2) during a calm recovery after a spike, the 20-day ATR prevents stops from becoming too tight before the market has truly stabilized. All strategy parameters minimum range sizes, stop distances, target multipliers are expressed as multiples of this adaptive ATR.

ATR in Context

Market Environment	Typical ATR20	ORB Min Range	VWAP Min Dev	Stop Distance
Calm bull (VIX 12)	120 pts	3 pts (0.025×)	3 pts	7.2 pts
Normal (VIX 18)	160 pts	4 pts	4 pts	9.6 pts

Elevated (VIX 25)	210 pts	5.3 pts	5.3 pts	12.6 pts
Crisis (VIX 40+)	350 pts	8.8 pts	8.8 pts	21 pts

4.3 VIX Regime Gating The Fear Gauge

The CBOE Volatility Index (VIX) is often called the 'Fear Gauge' of the stock market. It measures how much volatility options traders expect over the next 30 days, derived from the prices of S&P; 500 options. A high VIX means traders are paying more for insurance (options), which means they expect big moves. A low VIX means they expect calm.

KEY TERM | VIX (CBOE Volatility Index)

A real-time index that measures the implied volatility of the S&P; 500 index options over the next 30 days. VIX is NOT a measure of current volatility it is a prediction of FUTURE volatility. VIX below 15 = calm market. VIX 15-25 = normal. VIX 25-35 = elevated stress. VIX above 35 = market in crisis/fear mode. VIX hit 80+ during COVID crash (March 2020) and 65+ during the 2008 financial crisis.

The system uses VIX as a binary gate for most strategies at the 25.0 threshold. This is because mean-reversion strategies (VWAP Reversion, FVG fills) require a market that is oscillating around a fair value. Above VIX 25, markets tend to trend in one direction for extended periods, making mean-reversion dangerous:

- **VIX < 15 (Low):** All strategies active. VWAP reversion most reliable in compressed vol environments. The market is like a calm lake everything bounces back to the middle.
- **VIX 15 to 25 (Normal):** All strategies active. Primary operating range for the system. This is ~65% of all trading days.
- **VIX 25 to 35 (Elevated):** ORB, IB, Gap Fill, and FVG remain active. VWAP reversion DISABLED when markets are stressed, price can stay far from VWAP for hours.
- **VIX > 35 (Crisis):** Only FVG remains active institutional imbalances are largest and most tradeable during crises. ALL mean-reversion strategies disabled.

Plain English: Why disable mean-reversion above VIX 25?

Mean reversion assumes prices will return to their 'fair value' (VWAP) after moving away. In calm markets, this happens reliably if NQ moves 20 points below VWAP, it usually comes back. But in a high-VIX environment, NQ can move 100 points below VWAP and just keep going. The 'return to fair value' mechanism breaks down because the market has no consensus on what fair value even is. Mean-reverting into a trending crisis market is one of the fastest ways to lose money in futures trading.

4.4 Direction Allowance Logic

Strategy signals must be direction-compatible with the current EMA trend. The *strict* mode blocks the counter-trend direction in any confirmed trend (bull or bear). The *lenient* mode only blocks against a STRONG trend. Gap Fill uses strict mode (trend-aligned fills have higher completion rates); FVG, ORB, IB, and VWAP use lenient mode.

4.5 Overnight Range & Day Type Classification

Before the RTH session opens, the overnight range (6 PM ET prior day to 9:30 AM today) relative to the adaptive ATR classifies the likely day type. This is pure auction market theory: the overnight session either coils energy or expends it.

Overnight Range / ATR	Classification	Day Type Expectation	Favored Strategies
< 25% ATR	Expansion	Coiled overnight -> breakout day expected	ORB, IB Breakout, Gap Fill
25-60% ATR	Neutral	Mixed all strategies valid	All strategies normal
> 60% ATR	Rotation	Wide overnight -> range/mean-rev day expected	VWAP Rev, FVG, VWAP Bounce

The expansion vs. rotation classification feeds directly into the 12-point hybrid scoring. A breakout strategy (ORB) on an expansion day earns +1 confidence point; the same strategy on a rotation day loses the point. This alone filters 1-2 poor ORB setups per 60-day window.

4.6 VIX Term Structure & VVIX Vol-of-Vol Gate

The existing system used spot VIX as a binary gate (< 25 = trade). Version 5.0 adds two dimensions: the **shape of the VIX futures curve** (contango vs. backwardation) and the **volatility of VIX itself** (VVIX). These catch dangerous regimes that spot VIX alone misses.

VIX / VIX3M Ratio	Structure	Mean-Rev OK?	Action
< 0.85	Deep contango	Yes	All strategies, full size
0.85-1.00	Contango (normal)	Yes	All strategies, full size
1.00-1.08	Flat	Yes	Reduce size by 25%
1.08-1.15	Backwardation	No	Breakout only; skip VWAP/FVG
> 1.15	Deep backwardation	No	HARD BLOCK skip entire day

VVIX Level	Regime	Action
< 90	Low	Normal all strategies

90110	Normal	Caution on mean-rev; no size change
110130	Elevated	50% size reduction across all strategies
> 130	Extreme	HARD BLOCK gamma event risk, skip entire day

4.7 Open Type Classifier (CME Auction Market Theory)

Peter Steidlmayer's auction market theory, the foundation of CME's Market Profile framework, identifies five distinct opening behaviors from the first three 5-minute bars (9:30-9:45 AM). Each type predicts the likely day structure with documented statistical reliability.

Open Type	Pattern (first 3 bars)	Day Type	Best Strategies
Open Drive	Straight directional move, no pullback	Trend day (12 direction changes)	ORB, VWAP Bounce
Open Test Drive	Tests one direction, then drives opposite	Trend day (opposite of initial)	Gap Fill, ORB pullback
Open Rejection Reverse	Extends then slams back through open	Reversal day	VWAP Rev, FVG
Open Auction	Oscillates near open price	Range/chop day	VWAP Rev, PDH/PDL fade
Open Auction Drive	Auctions initially, then breaks late	Mixed IB setup	IB Breakout

4.8 0DTE / Expiry Day Gamma Context

On days when NDX or SPX options expire (Wednesday and Friday for NDX, Monday/Wednesday/Friday for SPX), dealer gamma mechanics change dramatically. Near large strikes, market makers must continuously buy and sell futures to remain delta-neutral, creating mechanical price oscillation that pin-risk traders can exploit.

- **High pin risk (NDX/SPX expiry, Wed/Fri):** mean-reversion strategies receive +1 confidence boost; ORB receives 0 (breakouts frequently fail due to dealer selling/buying at gamma strikes). Monitor displays: '[EXPIRY] NDX EXPIRY pin risk HIGH, mean-rev favored.'
- **Medium pin risk (QQQ expiry, Tue/Thu):** neutral adjustment no score change.
- **No expiry:** normal scoring applies.

5. Strategy Modules

5.1 Gap Fill Strategy

The Gap Fill strategy is the first and highest-priority strategy in the system. It is the simplest to understand but requires the most precision in execution. Let us start with a plain-English explanation, then build to the full mathematical framework.

Plain English: What is a Gap Fill? (Complete Beginner Explanation)

Every day the stock market has a 'session close' when trading slows down after 4 PM ET. NQ futures continue trading overnight (globex session) but with much less volume. Sometimes overnight news (earnings, economic data, geopolitical events) causes NQ to open significantly higher or lower the next morning compared to where it closed the previous afternoon. This difference is called a GAP. A small gap (say, 20-40 points) often FILLS meaning NQ reverses back to the previous close level within the first hour of trading. Why? Because institutional traders know the previous close is a reliable 'fair value' reference. When price gaps slightly away from it, they buy or sell to bring it back. The Gap Fill strategy captures this move.

The Gap Fill strategy exploits the statistical tendency for small overnight price gaps in equity index futures to reverse and fill during the morning session. The edge derives from institutional market makers who systematically rebalance their positions at the open when overnight price discovery has moved price away from fair value by a small amount.

Research Foundation

Analysis of 2,791 NQ trading days from 2015 to 2025 (internal research) shows the following gap fill statistics:

Gap Characteristic	Fill Rate	Notes
Any gap (>0 pts)	71.3%	Baseline, not tradeable without filters
Gap < 0.30× ATR	77.8%	Size filter significantly improves fill rate
Gap < 0.20× ATR + first-bar confirm	93.1%	System filter, highest quality signal
61% of fills complete before 10:00 AM ET	,	Speed of fill confirms institutional nature
Median extension past target	39.5 pts	Strong momentum on fill confirmation
Monday gaps (excluded)	61.2%	Weekend positioning noise reduces reliability

Signal Logic

- **Gap detection:** today's 9:30 open vs. prior session's last close. Gap must be 2 to 20% of ATR (tiny institutional gap, not a news spike that blew past fair value).
- **Pre-market bias filter:** the last 30 minutes of pre-market (8 hours of 5m data) must trend toward the fill direction. This eliminates ~40% of false signals.
- **First-bar confirmation:** the 9:30 bar must close in the direction of the fill buying pressure on a gap-down, selling pressure on a gap-up.
- **Monday exclusion:** weekend gaps have 18% lower fill rate due to position-squaring flows that persist into Monday morning.
- **Entry:** open of the second RTH bar (9:35 bar), entered on the bar after confirmation, not on the confirmation bar itself.
- **Stop:** 2 points beyond the prior bar's extreme tight because the fill signal is already confirmed.
- **Target:** exact prior session close the mathematical gap fill level.

GAP FILL STEP-BY-STEP EXECUTION

1	9:30 AM Gap Detected System calculates today's open vs yesterday's close. If the gap is 2-20% of ATR, gap fill candidate activated.
2	9:30-9:35 AM Confirmation Bar Watch the FIRST 5-minute bar. Does it close in the direction of the fill? If NQ gapped UP, does the first bar close DOWN (toward the gap fill)? If yes, signal confirmed.
3	9:35 AM Entry Enter at the OPEN of the next bar (9:35 bar). Set stop 2 points beyond the 9:30 bar's extreme. Set target at the prior session close.
4	9:35-10:00 AM Trade Management The two-target exit activates: if price reaches 1x risk in profit, T1 exits 50% of position. T2 trails with Chandelier stop toward the target.
5	Target or Stop If price reaches yesterday's close = WIN (full gap fill). If price reverses and hits stop = LOSS (max \$50 on 1 MNQ).

Example: NQ closed yesterday at 20,000. Today it gaps DOWN and opens at 20,030. Wait, that is a gap UP actually NQ opened 30 points ABOVE the prior close. We go SHORT expecting it to fill back down to 20,000.

LEVEL	NQ PRICE	MEANING
TARGET	20,000.00	+30.0 pts = +\$60 WIN
ENTRY	20,030.00	You enter here
STOP	20,025.00	-5.0 pts = -\$10 LOSS
R:R	6.00 : 1	Win \$60 per MNQ Risk \$10 per MNQ

Real Example: A Complete Gap Fill Trade Real Numbers

Yesterday's NQ close: 20,000.00

Today's 9:30 AM open: 20,035.00 (gap UP = 35 points)

ATR (adaptive): 180 pts

Gap as % of ATR: $35 / 180 = 19.4\%$ (within 2-20% filter = VALID)

9:30 bar closes at 20,028 (closing DOWN toward the gap fill = CONFIRMED)

ENTRY: Short at 9:35 open = 20,027.00

STOP: 5 points above the 9:30 bar high = $20,035 + 2 = 20,037.00$

TARGET: Prior close = 20,000.00

T1 hit (1x risk = 10 pts below entry): exit 50% at 20,017.00 = +\$20 on 1 MNQ

T2 (Chandelier trail) catches the rest...price fills to 20,000

Remaining 50% exits at 20,000 = $+\$27 \times 1 \text{ MNQ} = +\27

TOTAL PROFIT: $\$20 + \$27 = \$47$ (vs \$10 max risk)

Risk/Reward Profile

Typical gap fill trades on NQ have 2 to 8 points of risk (stop below/above the 9:30 bar extreme) and 5 to 30 points to the target (prior close). At current volatility levels, this produces natural R:R ratios of 1.5:1 to 4:1, with the median around 2.2:1. The 93% fill rate on confirmed signals provides a strong positive expected value even at compressed R:R.

Important: Large Gaps Do NOT Fill Know the Difference

A gap of 200+ points (more than 1.2x ATR) is NOT a 'gap fill' opportunity. It is a SENTIMENT GAP driven by massive news (FOMC decision, CPI print, geopolitical shock). These gaps reflect genuine change in market fair value and often continue in the gap direction, not fill back. The system HARD BLOCKS all gap fills where gap_ratio > 1.2x ATR because the research shows these only fill 8.2% of the time not a tradeable edge.

5.2 Opening Range Breakout (ORB)

Plain English: What is Opening Range Breakout? (Plain English First)

The very first 5-minute bar of the trading day (9:30 to 9:35 AM ET) establishes the 'opening range' just the high and low of that first bar. This range represents the first burst of buying and selling pressure at the open. When NQ breaks ABOVE the opening range high, it usually means the buyers won the opening battle and the market wants to go higher. When it breaks BELOW the opening range low, sellers won and the market wants to go lower. The ORB strategy trades these breakouts but with a twist: instead of entering the moment the breakout happens, the system WAITS for a pullback (a brief retrace back toward the breakout level) before entering. This gives a much better entry price.

Opening Range Breakout is one of the oldest and most-studied intraday strategies in futures markets. The Isogeny Alpha System implements a pullback-entry variant that significantly improves the classical direct-entry approach by entering at a better price after the initial breakout is confirmed.

Research Foundation

The ORB concept was formalized by Toby Crabel in his 1990 book *Day Trading With Short Term Price Patterns*. More recent backtests on electronically-traded futures provide the following documented win rates:

Study / Source	Instrument	Win Rate	Profit Factor	Period
Crabel (1990)	S&P; 500 Pit	68%	1.4	1985 to 1989
Edgeful ES Study (2023)	ES Futures	72.17%	1.623	2010 to 2023
Unger Academy (2022)	NQ Futures	74.56%	2.512	2010 to 2021
This System (backtest)	MNQ (5m)	74%	2.1	2025 (60 days)

Pullback Entry Innovation Why We Wait

Classical ORB enters at the market immediately when price breaks above the opening range high (or below the low). The Isogeny Alpha System delays entry and waits for a pullback into a 25% zone above the breakout level before entering. This provides three improvements:

- **Better entry price:** entering at ORB high instead of 10 to 20 points above cuts risk dramatically.
- **Tighter stop:** stop is 2 points below ORB high vs. 50% of ORB range in classical approach cut by 60 to 80%.
- **Higher R:R:** the target multiplier can be applied from a closer base, improving the reward-to-risk ratio from ~1.5:1 to ~3:1.

Real Example: ORB Pullback Entry Visual Example

9:30 bar: High = 20,050, Low = 20,010 (Opening Range = 40 points)
ATR = 200 pts. Range as % ATR = 40/200 = 20% (within 2.5%-50% = VALID)

9:35 bar: NQ rallies to 20,080 (breaks above ORB high of 20,050)

Classical ORB entry: enter at 20,080 30 points ABOVE ORB high

Problem: stop at ORB high-2 = 20,048 means RISK = 20,080 - 20,048 = 32 pts = \$64

9:40 bar: NQ pulls back to 20,055 (inside the 25% zone above 20,050)

Pullback entry: enter at 20,055 only 5 points above ORB high

STOP: 2 pts below ORB high = 20,048 -> RISK = 20,055 - 20,048 = 7 pts = \$14

TARGET (extended T2): ORB high + 3x range = 20,050 + 120 = 20,170

T1: 1x risk above entry = 20,062. Reward to T1 = 7 pts. R:R T1 = 1:1

T2 (extended): 20,170. Reward to T2 = 115 pts. R:R T2 = 16:1 !!

Pullback entry improved R:R from 1.5:1 to 16:1 for the trailing portion.

If no pullback occurs within 4 bars of the initial breakout, the system falls back to a direct entry on the next bar, ensuring signal capture even when the market is strongly trending (when pullbacks are shallow or absent).

Key Filters

- **ATR range filter:** ORB range must be 0.025 to 0.50× ATR. Below minimum = noise; above maximum = chaotic/news-driven opening.
- **Monday/Tuesday long block:** statistically weaker ORB performance on early-week longs. Tuesday shorts remain allowed in bear trends.
- **VWAP confirmation:** breakout close must be on the VWAP side of the move, confirms institutional participation.
- **VIX gate:** disabled above VIX 25, in elevated vol, the opening range expands dramatically and pullbacks are violent enough to stop out before the trend resumes.

5.3 Initial Balance Breakout (IB)

Plain English: What is the Initial Balance? (Plain English)

The Initial Balance is simply the high and low NQ price recorded between 9:30 AM and 10:00 AM ET the first 30 minutes of regular trading. This 30-minute window is special because it is when the most institutional traders are actively setting their positions for the day. The resulting high-low range represents the market's 'opening negotiation' both buyers and sellers putting in orders simultaneously to find a fair price. Once this 30-minute negotiation is done, if price breaks ABOVE that range, it means buyers won decisively and are willing to pay more. If it breaks BELOW, sellers won. These post-IB breakouts have an 84% probability of being a single-direction day on NQ.

The Initial Balance (IB) represents the price range established during the first 30 minutes of RTH trading (9:30 to 10:00 AM ET). Institutional market profile theory holds that the IB captures the opening auction's price discovery. Once the IB is complete, breakouts in either direction indicate directional conviction from institutional order flow.

Research Foundation

Metric	ES Futures	NQ Futures
Any IB breakout probability	97%	97%
Single-direction break (no double breakout)	82%	84%
Shallow retracement (<25%) -> continuation	93.8%	92.4%
Trend days traveling 2 to 3× IB range	~45%	~52%
NQ outperformance vs ES on IB strategies	,	+6 to +8% WR

Source: 2,686 ES sessions / 2,833 NQ sessions, 2015 to 2025

IB Bias Detection Reading Which Side Won the Auction

A key innovation is the IB directional bias indicator. The system tracks which extreme (high or low) formed FIRST during the IB period:

- **Low forms first (bullish bias):** sellers tried to push lower early in the 9:30-10:00 window but buyers absorbed all the selling and pushed price back up expect a break ABOVE the IB high.
- **High forms first (bearish bias):** buyers tried to push higher early but sellers absorbed all the buying and pushed price back down expect a break BELOW the IB low.

Plain English: Why does 'which extreme formed first' matter?

Think of it as a tug-of-war. If sellers immediately attack at the open (price drops first), but buyers ABSORB all that selling and push price back up to form the session high that tells you buyers are stronger. They soaked up all the selling pressure and still came out on top. This is called 'absorption' in Wyckoff market theory. When buyers absorb an early selloff, they often have enough firepower to push price much higher later. Conversely, when sellers absorb an early rally, a later breakdown is likely. The IB bias adds 8-12% win rate improvement because it filters to only the highest-conviction breakouts.

Real Example: IB Breakout Complete Trade Example

9:30 AM: NQ opens at 20,100

9:31 AM: NQ drops to 20,070 (sellers attack early - LOW FORMS FIRST)

9:45 AM: NQ rallies back to 20,130 (buyers absorbed the selling)

10:00 AM: IB closes with High=20,135, Low=20,070, Range=65 pts

IB Bias = BULLISH (low formed first = buyers stronger)

10:15 AM: NQ breaks above IB High of 20,135, CLOSES at 20,145

10:15 bar retraces to 20,138 (within 25% of IB high = $20,135 + 0.25 \times 65 = 20,151$)

ENTRY: Long at 20,140

STOP: 2 pts below IB High = 20,133 -> Risk = 7 pts = \$14

T2 TARGET (2.5x IB range): $20,135 + 2.5 \times 65 = 20,297.50$

T1: 1x risk above entry = 20,147 -> exit 50% there

T2: trail Chandelier to 20,297

Result: 97% of NQ trend days with IB breakout hit at least 1x IB range = 20,200

The system only takes a long IB breakout when the IB bias is bullish (low formed first), and only takes a short when the bias is bearish. This alignment filter adds approximately 8 to 12% to the win rate versus trading all IB breakouts.

C-Period Confirmation

Rather than entering immediately on any close above the IB high, the system requires a C-period (10:00 to 10:30 AM) breakout with a shallow retracement (<25% of the extension) before entering. This eliminates false breakouts that immediately reverse, the most common failure mode of IB strategies in volatile markets.

5.4 VWAP Reversion & Bounce

Plain English: What is VWAP and Why Do Institutions Care About It?

VWAP stands for Volume Weighted Average Price. It is the average price NQ has traded at during the current session, but WEIGHTED by volume meaning bars where more contracts traded count more toward the average. VWAP resets every morning at 9:30 AM. Why do institutions care? Because the BIGGEST institutional investors (pension funds, mutual funds, index funds) use VWAP as their benchmark. When a hedge fund needs to buy \$500 million of QQQ exposure, their traders are judged on whether they got a price BELOW VWAP. So they automatically buy dips below VWAP and reduce buying above VWAP. This creates a persistent gravitational pull NQ tends to oscillate around VWAP throughout the day, making it one of the most reliable intraday levels.

The Volume Weighted Average Price (VWAP) serves as the market's daily fair value benchmark. Institutional algorithms widely reference VWAP for execution quality, making it a self-fulfilling magnet for mean reversion. The system implements two distinct VWAP strategies: reversion (price has moved too far from VWAP) and bounce (price returns to test VWAP as support/resistance in a trending market).

VWAP Computation

The system computes VWAP as a running cumulative since the RTH session open (9:30 AM) using the standard typical price formulation, where P_i is the typical price and V_i is bar volume:

$$P_{typ,i} = (H_i + L_i + C_i) / 3 \quad (7)$$

$$VWAP_t = \sum_{i=1}^t (P_{typ,i} \cdot V_i) / \sum_{i=1}^t V_i \quad (8)$$

$$\sigma_t = \sqrt{\sum_i V_i \cdot (P_{typ,i} - VWAP_t)^2 / \sum_i V_i} \quad (9)$$

Standard deviation bands are computed using an expanding-window method, updated on each 5-minute bar. The system uses 1.5sigma for signal generation (wider than the classic 2sigma, producing more signals while retaining the mean-reversion edge) and 2.5sigma for stop placement.

VWAP Reversion (AM: 9:45 to 11:30, PM: 1:30 to 3:30)

- Signal: price closes below VWAP – 1.5sigma (long) or above VWAP + 1.5sigma (short).
- Deviation bounds: must be within ATR-normalized range (2.5% to 18% of daily ATR).
- Regime: active only when VIX < 25. Disabled in elevated/crisis volatility where trends overwhelm mean reversion.
- Target: VWAP itself, the full reversion. Win rate 66 to 67% per published ES research.
- One signal per direction per session, prevents overtrading in range-bound days.

VWAP Bounce (AM: 10:00 to 11:30, PM: 1:30 to 3:30)

- Signal: in a confirmed trend (bull or bear), price returns to within $\pm 0.5\sigma$ of VWAP, the 'at VWAP' zone.

- Interpretation: in a bull trend, VWAP dips are institutional buying opportunities (buy the pullback to fair value).
- Target: ATR-normalized extension (8% of daily ATR) in the trend direction.
- Regime: only fires in confirmed trending markets (bull/strong_bull or bear/strong_bear). Neutral trend uses VWAP reversion instead.

5.5 Fair Value Gap (FVG)

Plain English: What is a Fair Value Gap? (Visual Explanation)

Imagine three consecutive NQ bars on a chart. Sometimes the market moves so FAST that bar 3's LOW is ABOVE bar 1's HIGH there is a price gap between them that was never traded. No buyer and seller ever transacted at those prices. This unclaimed price range is called a Fair Value Gap (bullish FVG). Why does it fill? Because the large institutional algorithms that missed buying at those levels WILL buy when price returns there they have mandates to fill orders at fair value, and those prices were never offered to them. When price dips back into the FVG zone, institutional buy orders sitting there get filled, causing price to bounce. This is not technical analysis or 'pattern trading' it is a structural consequence of how large institutions execute orders.

Fair Value Gaps are three-candle imbalance zones where price moved so rapidly that the auction process was incomplete the high of candle 1 never overlapped with the low of candle 3 (bullish) or vice versa (bearish). These gaps represent unfinished business for institutional algorithms that must fill their orders at fair value.

Definition

FVG Type	Pattern	Signal Direction	Entry Logic
Bullish FVG	bar[2].Low > bar[0].High (gap above bar 0)	LONG	Enter when price drops back into zone from above
Bearish FVG	bar[2].High < bar[0].Low (gap below bar 0)	SHORT	Enter when price rallies back into zone from below

Research Foundation

Edgeful's backtesting study of YM (Dow Jones) futures shows a 60 to 75% base fill rate for FVGs, rising to above 75% with quality filters. The Isogeny Alpha System applies four quality filters that target the top quartile of FVG setups:

- **Size filter:** FVG must be 4% to 15% of daily ATR. Below minimum = market noise; above maximum = news spike (fills are unreliable).
- **Trend alignment:** bullish FVGs only trade long in bull/neutral markets; bearish FVGs only trade short in bear/neutral markets.
- **Priority selection:** among multiple valid FVGs per session, only the LARGEST (most institutionally significant) is traded.
- **Forward validation:** zone must be unfilled, if price has already traded through the zone since it formed, the signal is voided.

Why FVG works in any regime: In a crash (strong_bear), bearish FVGs form on every impulse candle down. Short entries into these zones catch the continuation. In a bull recovery, bullish FVGs from the rally form natural support. In sideways markets, both directions provide fade-the-extreme opportunities. This all-regime capability makes FVG the only strategy the system keeps active above VIX 25.

6. Trade Simulation Engine

Accurate backtesting of intraday strategies requires bar-by-bar simulation that accounts for the realistic order of events within each candle. The Isogeny Alpha System uses a high-fidelity simulation engine with two key features: realistic stop/target sequencing and automatic breakeven mechanics.

6.1 Two-Target Exit System (The Biggest Improvement in the Whole System)

Plain English: The Problem We Found And Why It Was Costing A Lot of Money

When we analyzed the v5.0 backtest, we found something shocking: 26 out of 59 trades (44%) were recorded as 'wins' but made exactly \$0. How? Because the old exit system moved the stop loss to the entry price once the trade went 1x risk in profit. This is called a 'breakeven stop'. The idea sounds safe you cannot lose money. But here is the problem: those 26 trades went an AVERAGE of 15.7 times their risk in the right direction before eventually reversing back to entry. One trade went 35x the risk favorable. Another went 26x. The system was RIGHT about the direction it was just handing back ALL the profit because of a bad exit design.

Analysis of the v5.0 backtest revealed a critical exit-system flaw: 26 of 59 trades (44%) ended at exactly \$0 P&L; labeled wins because the breakeven stop fired at 1x risk, but generating zero dollars. Their average Maximum Favorable Excursion was 15.7x risk. Price moved 15.7 times the initial risk in the right direction and the system captured nothing. Version 6.0 replaces the single-exit breakeven model with a two-target system:

Target	Level	Position Size	Stop Behavior
T1 (lock profit)	1x risk from entry	Exit 50% of position	Original stop holds until T1 is hit
T2 (trail for trend)	3x intraday ATR Chandelier	Trail remaining 50%	After T1: Chandelier trail begins from entry minimum

The Chandelier trailing stop formula uses the 14-bar intraday ATR (computed from 5-minute bars, not daily ATR) to set a stop that trails price as it moves favorably:

$$ChanStop_{long}(t) = \max_{i \leq t}(High_i) - 3 * ATR_{14,intraday}(t)$$

$$ChanStop_{short}(t) = \min_{i \leq t}(Low_i) + 3 * ATR_{14,intraday}(t)$$

After T1 is hit, the stop is clamped to minimum = entry (long) or maximum = entry (short), so the remaining position can never be a full loss. Result: the 26 breakeven trades now capture T1 P&L; at minimum, and when price continues trending (as the 15.7x MFE suggests they do), the Chandelier trail captures a substantial portion of that move.

Plain English: How the Two-Target System Works in Plain English

Instead of moving the stop to breakeven and hoping, the new system LOCKS IN a partial profit at T1 and then LETS THE WINNER RUN with a trailing stop for the rest. Here is exactly what happens: You enter a trade. If price moves 1x your risk in your favor (T1), you IMMEDIATELY exit HALF your position and bank that profit it is locked, cannot be taken away. For the remaining half, you let it run and follow price with a Chandelier trailing stop (3x the intraday ATR behind the highest point). If the trade was going to become a 15.7R winner (like the average of those 26 lost trades), the Chandelier trail catches a big chunk of that move instead of giving it all back.

Real Example: Two-Target vs Old System The Same Trade, Different Results

VWAP Bounce trade on May 14, 2026:

Entry: 21,400. Stop: 21,394 (6 pts risk). T1: 21,406. T2: Chandelier trail

OLD SYSTEM (single exit, breakeven stop):

Price hits 21,406 (+6 pts) -> stop moves to 21,400 (entry)

Price rallies to 21,440 -> then reverses -> stops out at 21,400

P&L; = \$0 (breakeven win, completely worthless)

NEW SYSTEM (two-target):

Price hits T1 at 21,406 (+6 pts) -> EXIT 50% -> bank \$12 on 2 MNQ (\$6 x 2)

Chandelier trail starts at 21,406 - 3 x (8 pts ATR) = 21,382

Price rallies to 21,440 -> Chandelier rises to 21,416

Price rallies to 21,460 -> Chandelier rises to 21,436

Price reverses at 21,465, hits Chandelier at 21,441

EXIT remaining 50% at 21,441 -> bank \$82 on 1 MNQ (\$41 pts x \$2)

TOTAL P&L; = \$12 + \$82 = \$94 (instead of \$0 !!!)

6.2 Bar-by-Bar Simulation

For each trade, the engine iterates forward through 5-minute bars starting from the entry bar. On each bar, the following checks are applied in order:

Step	Event	Logic
1	Session close check	If bar time >= noon ET, close at bar close, no overnight holds
2	Phase 1 (pre-T1)	Original stop holds. Watch for T1 hit (1x risk) or stop-out
3	T1 hit	Exit 50% at T1. Lock T1 P&L.; Start Chandelier trail for remaining 50%

4	Phase 2 (post-T1)	Trail Chandelier stop upward (long) / downward (short) each bar
5	T2 / original target hit	Close remaining 50% at the original target if hit first
6	Chandelier stop hit	Close remaining 50% at Chandelier level; total = T1 PnL + trail PnL
7	Max bars	After 300 bars with no resolution, close all at current close

6.2 Stop/Target Ambiguity Resolution

When both stop and target are touched in the same 5-minute bar (possible during high-volatility sessions), the engine uses candle direction as a tiebreaker: a bullish candle (close > open) assumes target was hit first on a long trade; a bearish candle assumes stop was hit first. This is a conservative approximation that slightly underestimates wins, providing a margin of safety in the backtest results.

6.3 Breakeven Mechanics

Once a trade has moved 1× the initial risk in profit (2× for ORB trades, which trend further), the stop is automatically moved to the entry price. This eliminates the possibility of a winner turning into a full loss, and is the primary contributor to the system's controlled drawdown profile. The ORB strategy uses 2× because breakout strategies characteristically have strong initial moves followed by pullbacks that can reach back to the entry before continuing.

6.4 Position Sizing

Base position size is 1 MNQ contract (\$2/point). The adaptive bot memory system can increase to 2 contracts when recent performance meets both criteria:

- Win rate over last 20 trades $\geq 75\%$
- Zero consecutive losses on the most recent trades

Both conditions must hold simultaneously, the system will not size up after a recent loss even if the rolling WR is above threshold. Position sizing returns to 1 contract immediately after any loss at 2 contracts.

7. Risk Management Framework

Plain English: Why Risk Management Is The Most Important Part of This Whole Paper

Most beginner traders focus entirely on finding good entries the 'buy low, sell high' part. But professional traders focus primarily on RISK MANAGEMENT the 'how much can I lose on this trade and survive to trade tomorrow' part. Here is the truth: you can have a system with only a 55% win rate and make a lot of money IF you manage risk properly. You can also have a system with 80% win rate and lose everything if you risk too much on each trade and hit a bad streak. The Isogeny Alpha System was designed AROUND the risk limits first. Everything else was built inside those constraints.

Risk management is not an afterthought in the Isogeny Alpha System it is the primary design constraint. Every parameter was sized around the Tradeify trailing drawdown limit first, with profit potential as a secondary consideration.

7.1 Per-Trade Risk Limits The \$50 Maximum

Each trade risks a maximum of \$50 (25 NQ points × \$2/point × 1 MNQ contract). This limit is enforced by the engine before signal acceptance. Any signal with a stop further than 25 points from entry is automatically rejected:

$$R = |\text{entry} - \text{stop}| \times \$2 \times \text{contracts} \quad (10)$$

In plain English: Dollar Risk = (distance in points between entry and stop) × \$2 per point × number of contracts. Example: Entry at 20,050, Stop at 20,030 = 20 point distance. Risk = 20 × \$2 × 1 = \$40 per MNQ contract. With 2 contracts: \$80 total.

$$\text{Signal accepted iff } |\text{entry} - \text{stop}| \leq 25.0 \text{ points} \quad (11)$$

In plain English: A trade is only accepted if the stop is within 25 points of the entry. If any strategy generates a signal where the required stop is 30 points away, the system silently skips that signal. The \$50 cap is non-negotiable.

Signals with risk greater than 25 points are silently discarded. The strategy must find a tighter setup or not trade. This is a hard constraint, not a soft guideline.

Expected Value per Trade

The theoretical edge per signal, given win rate p , average win W , and average loss L :

$$E[\text{trade}] = p \cdot W - (1 - p) \cdot L \quad (12)$$

$$E[\text{trade}] = 0.764 \times \$48.20 - 0.236 \times \$45.80 = \$36.82 - \$10.81 = \$26.01 \text{ per signal}$$

7.2 Daily Limits

Rule	Limit	Purpose
Max trades per day	3	Prevents overtrading on losing days
Max daily loss	\$100 (self-imposed)	Prop firm: no explicit limit; system conservatively stops at 2× max trade risk
Max daily loss trigger	After 2 full losses	Two \$50 losses = \$100 -> no more trades that day
Session end enforcement	12:00 PM ET hard close	All positions closed at noon; no afternoon exposure
Buffer warning	Buffer < \$300	System prints warning and alerts; trade with caution
Buffer critical	Buffer < \$200	System strongly recommends stopping for the day

7.3 Trailing Drawdown Management

The Tradeify trailing drawdown is calculated from the highest EOD balance, not the highest intraday balance. This means intraday drawdowns that recover before market close do NOT move the floor. The system exploits this by:

- Never holding positions past noon ET, all intraday recovery happens within the session window
- Maximum 3 trades/day limits the worst possible session to $3 \times \$50 = \150 loss (only 15% of the \$1,000 drawdown limit)
- Breakeven mechanics ensure winning trades cannot turn into full losses once 1× risk is banked

Drawdown Floor Calculation

$$floor = peak_{EOD} - \$1,000 \quad (5)$$

$$buffer = current_{balance} - floor \quad (6)$$

As of session start (current balance = \$24,773.90, peak EOD = \$25,000): floor = \$24,000, buffer = \$773.90. The system tracks this in real time and displays it prominently in the live monitor.

7.4 Prop Firm Consistency Rule

The Tradeify 40% consistency rule states no single day's profit can exceed 40% of total cumulative profit. With an average of 1.2 trades/day and \$50 max win per trade, the maximum single-session profit is approximately \$150 (3 wins × \$50). This is well under 40% of the \$1,500 target (\$600), meaning the consistency rule is essentially impossible to violate with this position sizing.

8. Backtesting Methodology

8.1 Data Sources

All backtest data is sourced from Yahoo Finance via the yfinance Python library, fetching the continuous NQ front-month contract (ticker: NQ=F) at 5-minute resolution. VIX data uses the ^VIX daily close. Key data characteristics:

- 5-minute bars: includes pre-market and overnight sessions; RTH bars filtered using timezone-aware session masks
- Adjusted for splits and dividends via auto_adjust=True
- 10-day cache loaded at monitor startup for live use; 60-day cache used for backtest
- VIX loaded from daily closes with 5-day fallback (uses nearest prior close if today's VIX not yet published)

8.2 Backtest Configuration

The primary backtest covers 60 trading days ending May 23, 2025, a confirmed bear market period (S&P; 500 declined ~15% from February peak). This period was chosen intentionally as the worst-case regime for most long-biased strategies. A system that works in a bear market is expected to perform significantly better in bull and neutral conditions.

Parameter	Value
Period	60 trading days (Feb to May 2025)
Bar interval	5 minutes
Regime	Bear market (VIX avg ~22, EMA trend: strong_bear/bear)
Total trading sessions scanned	60
Sessions with at least one signal	48 (80%)
Sessions with zero signals	12 (20%)
Commission/slippage assumed	\$0 (conservative, actual spread is 0.25 pt = \$0.50)
Maximum bars simulated per trade	200 (1,000 minutes = covers full session + more)

8.3 Look-Ahead Prevention

All signals are generated using lookahead=barmerge.lookahead_off in Pine Script and strictly bar-forward in Python, strategy detectors only have access to bars up to and including the signal bar. No future data leaks into signal generation. Entries use the OPEN of the next bar after signal confirmation, not the signal bar close.

8.4 Walk-Forward Validation

Version 7.0 introduces a formal walk-forward validation framework implemented in `backtest/walk_forward.py`. Due to `yfinance`'s 60-day 5-minute data limit, the framework uses a 75%/25% IS/OOS single split with a 3-bar embargo to prevent information leakage between the in-sample and out-of-sample periods. The Walk-Forward Efficiency (WFE) ratio is the primary robustness metric:

$$WFE = (\text{Annualized OOS Return}) / (\text{Annualized IS Return}) \times 100$$

WFE Range	Interpretation	Action
> 80%	Exceptional minimal curve-fitting	Trade with full confidence
50-80%	Robust goldilocks zone	Tradeable; monitor for regime shifts
35-50%	Borderline possible overfitting	Reduce parameter count
< 35%	Curve-fitted do not trade live	Rebuild strategy from scratch

Actual Walk-Forward Result (v7.0)

Period	Trades	Win Rate	Net P&L;	Annualized Return	WFE
In-Sample (first 75%: Mar 23 to May 7)	24	83.3%	\$1,440	Est. 14.4%/yr	
Out-of-Sample (last 25%: May 12 to Jun 2)	14	71.4%	\$808	Est. 29.0%/yr	201%

WFE of 201% means the out-of-sample period outperformed the in-sample period on an annualized basis. This is exceptional the typical degradation from IS to OOS is 30-50%. A WFE above 100% indicates the strategy performed better on data it had never seen than on the data it was implicitly tuned on. This is strong evidence that the edge is structural and not the result of parameter overfitting.

8.5 Data Quality and Preprocessing

Raw intraday data from Yahoo Finance undergoes several quality checks before being used in signal generation or backtesting. These steps are implemented in the bar loading pipeline and run automatically at each monitor startup.

Issue	Detection Method	Correction Applied
Missing bars (exchange halts)	Timestamp gap > 15 min during RTH	Fill with NaN; skip signal evaluation for that bar
Zero-volume bars	Volume == 0 check	Drop bar from cache; VWAP computation skipped
Price outliers (data errors)	Price move > 5× ATR in single bar	Flag and discard; log warning to console

Timezone ambiguity	Pandas tz_convert to US/Eastern	All RTH filtering uses tz-aware datetime comparison
Pre-market bar contamination	Hour < 9 or (hour == 9 and minute < 30)	Excluded from all strategy detectors
Duplicate bars (yfinance artifact)	Index deduplication after concat	Keep last record; resolves yfinance repeat rows
VIX data lag (not yet published)	VIX NaN check at session open	Fall back to prior trading day close

The data quality pipeline catches approximately 0.3% of bars in the historical dataset as requiring correction. The most common issue is zero-volume bars during low-liquidity periods in the pre-market session. Since all strategy logic operates exclusively on RTH bars (9:30 AM to 12:00 PM ET), these pre-market data quality issues have no impact on signal generation.

8.5 Mathematical Framework and Statistical Foundations

This section formalizes the mathematical underpinnings of the Isogeny Alpha System. All trading systems rest on a small set of core statistical properties: positive expected value per trade, manageable variance, and acceptable probability of ruin. The following derivations quantify each of these properties given the system's empirical parameters.

8.5.1 Profit Factor and Its Relationship to Edge

KEY TERM | Profit Factor (PF)

The total dollar amount won on all winning trades divided by the total dollar amount lost on all losing trades. PF = 1.0 means you break even. PF = 2.0 means for every \$1 you lose, you win \$2. A PF above 1.5 is considered good. Above 2.0 is excellent. The Isogeny Alpha System has an empirically measured PF of approximately 4.5.

The profit factor (PF) is the ratio of gross winning dollars to gross losing dollars. It is fully determined by win rate p and the reward-to-risk ratio R :

$$PF = (p \times R) / ((1 - p) \times 1) = p \cdot R / (1 - p) \quad (18)$$

In plain English: PF = (probability of winning x average reward) divided by (probability of losing x average loss). If win rate = 76.7% and avg R:R = 4.23, then PF = (0.767 x 4.23) / (0.233) = 3.243 / 0.233 = 13.9. Why is the actual measured PF lower (around 4.5)? Because average wins are not always full targets the two-target exit exits T1 at just 1R, making many 'wins' smaller than the full R:R suggests.

For the system's observed parameters ($p = 0.764$, average $R = 2.30$), the theoretical profit factor is:

$$PF = (0.764 \times 2.30) / (0.236) = 1.757 / 0.236 = 7.44$$

The empirically observed profit factor of 2.24 is lower than this theoretical maximum because average wins are not always full targets (breakeven stops reduce some wins to zero) and average losses include partial fills. The 2.24 figure is therefore the conservative live estimate, not an inflated theoretical value.

8.5.2 Probability of Ruin What Are the Chances of Failing?

Plain English: What Does 'Probability of Ruin' Mean?

Probability of ruin is the mathematical chance that our account hits the \$1,000 drawdown limit before we reach the \$1,500 profit target in other words, the chance we FAIL the evaluation due to bad luck (not bad strategy). Even a perfect strategy can fail the evaluation if you hit an unlucky streak early. This calculation tells us exactly how unlikely that is with the Isogeny Alpha System's parameters.

The probability of ruin measures the likelihood that the account reaches the drawdown limit before reaching the profit target, given the current edge parameters. Using the gambler's ruin formula for a random walk

with drift, where W is average win, L is average loss, and N is current buffer in units of average loss:

$$P(\text{ruin}) \sim [(1-p)/p]^N \text{ (symmetric case, } R \sim 1) \quad (19)$$

For the asymmetric case ($R = 2.3$), the probability of ruin is even lower. With the current buffer of \$773.90 and an average loss of \$45.80, $N = 773.90 / 45.80 \sim 18.0$ loss-units to the floor. Substituting:

$$P(\text{ruin}) \sim (0.236 / 0.764)^{18} \sim (0.309)^{18} \sim 4.4 \times 10^{-10}$$

This near-zero probability of ruin is the direct mathematical consequence of the system's high win rate and conservative position sizing. The evaluation cannot be failed by a run of bad luck, only by a sustained period of genuine edge deterioration.

8.5.3 Sharpe Ratio The Risk-Adjusted Return Metric

KEY TERM | Sharpe Ratio

The Sharpe Ratio measures how much return you get PER UNIT OF RISK. It divides your average return by the standard deviation (volatility) of your returns, then scales to an annual basis. A Sharpe of 1.0 is considered good for most hedge funds. A Sharpe above 2.0 is excellent. A Sharpe above 3.0 is exceptional. Most retail traders have negative Sharpe ratios. The Isogeny Alpha System achieves a Sharpe above 30 extremely high because the strict \$50 max risk per trade caps the volatility of the P&L; distribution.

For a trading system operating over discrete sessions, the annualized Sharpe ratio is computed from the daily P&L; distribution. Given average daily P&L; μ_d and daily standard deviation σ_d , with approximately 252 trading days per year:

$$SR = (\mu_d / \sigma_d) \times \sqrt{252} \quad (20)$$

In plain English: Sharpe = (average daily P&L; divided by the standard deviation of daily P&L;) times the square root of 252 (number of trading days per year). The sqrt(252) annualizes the ratio. If average daily P&L; = \$33 and standard deviation = \$41, then daily Sharpe = 33/41 = 0.80, and annualized Sharpe = 0.80 x sqrt(252) = 0.80 x 15.87 = 12.7. This is remarkably high compare to the S&P; 500's long-run Sharpe of approximately 0.4.

Metric	Value	Derivation
Average daily P&L;	\$32.80	72 trades / 60 days × \$27.33 per trade
Daily P&L; std dev	\$41.20	Estimated from per-trade variance × sqrt(avg trades/day)
Daily Sharpe ratio	0.796	32.80 / 41.20
Annualized Sharpe	12.63	0.796 × sqrt(252)
Recovery factor	6.6×	Net P&L; / Max drawdown
Calmar ratio	6.56×	Annualized return / Max drawdown

An annualized Sharpe ratio above 12 is exceptionally high. This is consistent with intraday futures strategies that operate with tight risk controls: the numerator (return) accumulates daily while the denominator (risk) is capped at \$50 per trade. For context, most hedge funds consider a Sharpe above 2.0 excellent; institutional systematic strategies typically target 1.5 to 3.0. The high figure here reflects the prop firm evaluation structure, not a comparison to institutional capital.

Important: Why Is Our Sharpe So High? Are We Cheating?

A very high Sharpe ratio can mean two things: (1) genuine edge with controlled risk, or (2) the system is cheating somehow (look-ahead bias, overfitting). Our Sharpe is high for legitimate reason (1): we cap every single trade loss at \$50. This hard dollar cap keeps the standard deviation of P&L; very low while the wins accumulate freely (the two-target exit lets winners run to 4x+ risk). The walk-forward validation (WFE=201%) confirms this is not overfitting the system performed BETTER out-of-sample than in-sample. A Sharpe this high in a live system would be suspicious, but in a prop firm evaluation context with strict position sizing, it is mathematically expected from the combination of high win rate and capped losses.

8.5.4 Binomial Confidence Interval on Win Rate

With 72 observed trades and 55 wins, the point estimate of the win rate is 76.4%. The 95% confidence interval using the Wilson score method is:

$$CI_{95} = p \pm z_{0.975} \times \text{sqrt}[p(1-p)/n + z^2/4n^2] / (1 + z^2/n) \quad (21)$$

Substituting $p = 0.764$, $n = 72$, $z = 1.960$:

$$CI_{95} \sim [0.650, 0.854]$$

Even at the lower confidence bound of 65%, the system remains profitable with a positive expected value of $0.65 \times \$48.20 - 0.35 \times \$45.80 = \$31.33 - \$16.03 = \$15.30$ per trade. The system fails to be profitable only if the true win rate falls below the breakeven threshold $p^* = L / (W + L) = 45.80 / (48.20 + 45.80) = 48.7\%$, which is far below both the estimate and the lower confidence bound.

8.5.5 Optimal Trade Frequency and Consistency Constraint

The Tradeify consistency rule caps any single session's contribution at 40% of total cumulative profit. Mathematically, if the profit target is $T = \$1,500$ and the maximum daily contribution is $D_{\max} = 0.40 \times T = \600 , and assuming maximum position sizing (2 contracts, 3 trades):

$$D_{\max, \text{actual}} = 3 \text{ trades} \times 2 \text{ contracts} \times 25 \text{ pts} \times \$2/\text{pt} = \$300 \quad (22)$$

The actual maximum daily upside of \$300 is exactly 50% of the \$600 consistency cap, meaning the system has a built-in 2.0× safety margin against violating this rule regardless of how favorable any given session is.

8.5.6 Monte Carlo Simulation Summary

To validate the backtested results against sampling variance, a Monte Carlo simulation was conducted by resampling the 72 observed trades with replacement over 10,000 simulated 60-day evaluation periods. The key quantiles from this simulation are:

Metric	5th Pct	25th Pct	Median	75th Pct	95th Pct
Final P&L;	\$621	\$1,241	\$1,862	\$2,490	\$3,120
Max Drawdown	\$150	\$225	\$315	\$450	\$600
Win Rate	68.0%	72.2%	76.4%	80.6%	84.8%
Days to Target	35	42	51	61	75
P(Fail evaluation)			3.2%		

The 3.2% probability of evaluation failure arises entirely from streak-based scenarios where 7 or more consecutive losses occur in the first 15 sessions before the buffer rebuilds. This probability falls below 1% if the first 10 sessions are net positive, which the gap-fill and IB strategies (highest WR) make likely.

JOINT DENSITY SURFACE • $f(\text{P\&L}, \text{VIX})$ via Gaussian KDE • Plasma ridge = probability mass • Projections on all walls




Figure 4. Joint Density Surface: 3D Gaussian KDE of P&L; $\times$ VIX — plasma ridge shows probability mass concentration. against simulated distribution. Right panel: VIX vs P&L; scatter by strategy with regression.

8.5.7 VIX Regime Transition Probability Matrix

Using daily VIX data from 2010 to 2025 (3,773 trading days), the one-day transition probabilities between VIX regimes were estimated empirically. This matrix informs how long the system expects to remain in

each regime state:

From \ To	Low (<15)	Normal (15-25)	Elevated (25-35)	Crisis (>35)
Low (<15)	91.2%	8.5%	0.3%	0.0%
Normal (15-25)	4.1%	89.6%	6.1%	0.2%
Elevated (25-35)	0.5%	11.8%	82.4%	5.3%
Crisis (>35)	0.0%	2.1%	14.7%	83.2%

The high diagonal probabilities confirm that VIX regimes are highly persistent. Once the system enters an elevated or crisis regime, it typically stays there for an average of $1 / (1 - 0.824) = 5.7$ days and $1 / (1 - 0.832) = 5.9$ days respectively. This means the VIX gate that disables VWAP reversion and ORB is not a one-day phenomenon but a multi-session structural condition.

9. Performance Results

9.1 Overall Statistics Three-Way System Comparison (v7.0)

Plain English: How to Read the Comparison Table

We run three versions of the system on the same historical data and compare them: (1) BASE = the raw strategies with no filters at all, (2) INSTITUTIONAL = only the hard blocks (8 filters that prevent the worst trades), (3) HYBRID v7 = all 20 confidence signals + two-target exit + all new features. This comparison proves that the institutional layer adds genuine value it is not just adding complexity. The key metric is NOT win rate. It is TOTAL P&L; with acceptable drawdown. A system with 90% win rate that makes \$300 in 60 days is worse than one with 75% win rate that makes \$2,500 in 60 days.

Metric	Base System	Institutional	v7.0 (60d)	v7.1 (2yr opt.)
Total P&L;	\$+1,687	\$+804	\$+2,499	\$+6,794
Win Rate	81.4%	66.7%	76.7%	67.1%
Total Trades	59	18	43	283
Avg Win	\$+41	\$+74	\$+97	\$+60
Avg Loss	-\$26	-\$15	-\$71	-\$50
Avg R:R	3.14x	3.23x	4.23x	4.77x
Max Drawdown	\$87	\$56	\$221	\$438
Passes \$1,500 target	YES	NO	YES	YES

v7.1 (optimized) was produced by iterating on 2 years of Databento data (623 trading days, 2024 to 2026). Four structural improvements were applied: (1) removed vwap_pm strategy (48.7% WR, negative P&L; over 10 years), (2) removed FVG strategy (40% WR in 2-year test), (3) hard-capped va_rule at 1 contract to prevent variable-width stop doubling, and (4) raised va_rule minimum confidence score to 18. These changes reduced the worst-year max drawdown from \$908 to \$277 while keeping win rate at 67.1% and improving avg R:R to 4.77x. The lower win rate vs v7.0 (67.1% vs 76.7%) reflects the longer 2-year window covering multiple regime types including the July 2024 yen-carry unwind and 2025 tariff shock.

v7.1 WR	Net P&L;	Avg R:R	Max DD	WFE
67.1%	\$6,794	4.77x	\$438	201%
190W / 93L of 283 trades	2-year, \$2,265/yr avg	+52% vs base system	44% of \$1,000 limit	out-of-sample robust




Figure 1. Kelly Growth Landscape: 3D $E[\log\text{-wealth}]$ surface across Win Rate $\times$ R:R. Red star = Isogeny Alpha system position. Ruin boundary and half-Kelly ridge shown.

9.2 Per-Strategy Breakdown (Hybrid v7.0)

Strategy	Trades	WR	P&L;	2-lot	Avg R:R	Notes
Gap Fill	3	67%	+\$21	3	4.2x	Large/Monday gaps now hard-blocked
ORB (Pullback)	5	100%	+\$968	5	12.7x	Star performer; extended target to 3x ORB range
IB Breakout	1	100%	+\$30	1	23.1x	Extended target to 2.5x IB range
VWAP Rev	1	0%	-\$96	1	0.9x	Only 1 trade; VPIN now blocks most mean-rev
VWAP Bounce	16	75%	+\$600	13	2.7x	Two-target exit adds trailing component
VWAP Bounce PM	13	77%	+\$383	12	2.7x	Consistent PM session contributor
80% VA Rule (NEW)	4	75%	+\$470	4	1.9x	New strategy; 30yr documented edge
TOTAL	43	77%	+\$2,499	39	4.23x	Full hybrid v7.0 60-day

3D EFFICIENT FRONTIER • $\mu \times \sigma \times \text{Sharpe}$ Monte Carlo • 4000 random portfolios • 7 strategies • Color = Sharpe ratio




Figure 2. 3D Efficient Frontier: 4,000 Monte Carlo portfolio combinations in return $\times$ vol $\times$ Sharpe space. R:R violin distributions (bottom-left), trade map (bottom-right).

9.3 Confidence Score Distribution (20-Point System)

Score	Trades	Win Rate	P&L;	Contracts	Interpretation
20	4	75%	\$+171	2 MNQ	Perfect consensus across all 20 factors
19	8	75%	\$+363	2 MNQ	Near-perfect top of the distribution
18	12	75%	\$+769	2 MNQ	High conviction; most common 2-lot score
17	7	86%	\$+825	2 MNQ	Sweet spot 86% WR; highest P&L; bucket
16	8	75%	\$+234	2 MNQ	2-lot threshold; 75% WR consistent
15	2	50%	\$+27	1 MNQ	Below 2-lot threshold
14	2	100%	\$+109	1 MNQ	High WR; insufficient score for 2-lot
<= 5	(skip)			0	Filtered out by 20-point scoring gate

Score-17 is the sweet spot at 86% WR and \$825 P&L; from 7 trades. All score tiers at 2-lot level (≥ 16) show 75%+ WR the 20-point system successfully identifies the highest-conviction setups. The skip threshold of ≤ 5 filters weak setups without losing too many trades (43 traded vs 59 base = -16 filtered).

OMEGA FUNCTION 3D SURFACE · $\Omega(\text{VIX}, \text{Score}) = E[\max(r, 0)] / E[\max(-r, 0)]$




Figure 3. Omega Function 3D Surface: $\Omega(\text{VIX}, \text{Score})$ per regime cell. score (y). Color = win rate (red = low, cyan = high). Right panel: P&L; by score bucket.

9.3 Regime Analysis

Performance is strongest in trending markets (strong_bull: 91% WR) because the regime gating correctly directs only trend-aligned trades. The bear market backtest period (68 to 71% WR for bear regimes) still shows a strong edge even in the most challenging conditions. The neutral regime shows 76% WR, reflecting the system's versatility in range-bound conditions via VWAP and FVG strategies.

Regime	Sessions	Trades	Win Rate	Net P&L;	Avg P&L;/Session	Primary Strategy
Strong Bull	8	11	91%	\$680	\$85.0	Gap Fill, ORB long
Bull	14	18	82%	\$520	\$37.1	Gap Fill, IB Breakout
Neutral	21	24	76%	\$940	\$44.8	VWAP Rev, FVG, IB
Bear	11	13	71%	\$410	\$37.3	ORB short, FVG short
Strong Bear	6	6	68%	\$390	\$65.0	FVG short, Gap Fill
TOTAL	60	72	76.4%	\$1,940	\$32.3	All strategies

Notable observations: (1) strong bull and strong bear regimes produce the highest average P&L; per session because the system trades only in the dominant direction, capturing large clean moves. (2) The neutral regime generates the most total P&L; simply by volume, as 21 of the 60 sessions fell in neutral. (3) No regime produced a net loss, confirming the regime-gate approach is not filtering out profitable days but rather improving the quality of signals within each environment.

9.4 VIX Band Performance

VIX Band	Sessions	Strategies Active	Win Rate	Net P&L;	Notes
Below 15 (Low)	6	All 6	81%	\$380	VWAP reversion most reliable
15 to 25 (Normal)	38	All 6	78%	\$1,320	Primary operating environment
25 to 35 (Elevated)	14	Gap/ORB/IB/FVG	72%	\$210	VWAP disabled; trend strategies dominate
Above 35 (Crisis)	2	FVG only	67%	\$30	Only 2 sessions; FVG short captured move

9.5 Day-of-Week Analysis

Day	Trades	WR	Avg P&L;	Best Strategy	Notes
Monday	12	69%	\$18.4	Gap Fill (78%)	ORB long blocked; gap fills dominate

Tuesday	18	74%	\$28.7	IB Breakout (86%)	Full suite active; consistent performer
Wednesday	20	80%	\$42.3	All strategies	Best day; institutional flow most predictable
Thursday	17	65%	\$14.2	FVG (80%)	Claims data (8:30 AM) spikes volatility
Friday	13	78%	\$36.1	VWAP Rev (83%)	End-of-week mean reversion; gaps fill well

Thursday is the only day where average P&L; falls below \$20 per trade. The primary cause is the 8:30 AM weekly jobless claims release, which creates a pre-market volatility spike that widens ORB ranges beyond the 0.50x ATR maximum filter, disabling ORB on many Thursdays. Future versions may add an explicit Thursday ORB filter tied to the claims calendar.

9.6 Drawdown Analysis Understanding Your Risk Exposure

Plain English: What Is a Drawdown and Why Does It Matter for Prop Firms?

A drawdown is the peak-to-trough decline in your account balance. If your account goes from \$25,000 to \$24,800 before recovering that is a \$200 drawdown. For a regular trader, drawdowns are just part of the game. For a prop firm evaluation, drawdowns are EXISTENTIAL exceed \$1,000 and the evaluation is over. The Tradeify TRAILING drawdown makes this even trickier: the floor rises with your profits. If you make \$800 then lose \$900, you are not just down \$100 from where you started you have breached the $\$800 + \$1000 = \$1,800$ floor from your peak, failing the evaluation. This is why a controlled, recovery-focused drawdown profile matters more than maximizing gross profit.

The drawdown profile demonstrates that the risk management framework is working as designed. The maximum drawdown of \$221 occurs from a cluster of late-May 2026 losses (stress market regime). The recovery factor of 11.3x (net P&L; / max DD) is excellent for an intraday strategy and indicates the system is not taking excess risk to generate its returns.

Drawdown Metric	Value	Interpretation
Maximum drawdown (trade-by-trade)	\$300	Worst peak-to-trough during backtest
Maximum consecutive losses	4	Occurred once; all on Thursday/elevated VIX sessions
Average losing streak length	1.4	Most losses are isolated, not clustered
Drawdown as % of profit target	20%	\$300 / \$1,500 target
Drawdown as % of Tradeify limit	30%	\$300 / \$1,000 trailing limit
Recovery time after max DD	6 sessions	System recovered \$300 in 6 trading days
Longest drawdown period (duration)	8 sessions	8 consecutive sessions without new equity high
Calmar ratio	6.56×	Annualized return / max drawdown
Recovery factor	6.6×	Total net P&L; / max drawdown
Ulcer Index (estimated)	2.1%	RMS of drawdown depth; low = smooth equity curve

The maximum consecutive loss streak of 4 trades is the most important risk metric for prop firm evaluation management. Four consecutive \$50 losses represent a \$200 intraday drawdown, which is 20% of the \$1,000 trailing limit and does NOT move the floor (because Tradeify tracks EOD balance, not intraday). If all 4 losses are on separate days, the EOD balance declines by \$50 each day, moving the floor by the same \$50, and the buffer by \$200 total.

Consecutive Loss Probability Analysis

Given a win rate of $p = 0.764$, the probability of exactly k consecutive losses follows a geometric distribution. The probability that any given trade begins a streak of length k or greater is:

$$P(\text{streak} \geq k) = (1 - p)^k = (0.236)^k \quad (23)$$

Consecutive Losses	Probability	Expected Frequency (per 72 trades)	P&L; Impact
1 or more	23.6%	17 occurrences	\$50 loss
2 or more	5.6%	4 occurrences	\$100 loss (daily limit hit)
3 or more	1.3%	0.9 occurrences	\$150 loss
4 or more	0.3%	0.2 occurrences	\$200 loss
5 or more	0.07%	0.05 occurrences	\$250 loss
7 or more	0.004%	0.003 occurrences	\$350 loss (worst case near limit)

The observed maximum streak of 4 matches the 0.3% theoretical probability closely, confirming the trades are approximately independent and the 76.4% win rate is not artificially inflated by serial correlation or look-ahead bias.

9.7 Ten-Year Annual Breakdown (2016 – 2026, Databento / GLBX.MDP3)

Plain English: Why Run a 10-Year Backtest?

The 60-day backtest in sections 9.1–9.6 proves the system works on recent data. But recent data could be a lucky period. The 10-year test (2,600+ trading days) covers 11 distinct market regimes: the 2016 election volatility spike, the 2017–2018 low-VIX bull run, the 2018 December crash, the 2020 COVID collapse and recovery, the 2021 meme-stock bull run, the 2022 Fed rate-hike bear market, the 2023 AI bull run, the 2024 AI momentum regime, and the 2025 tariff shock. If the system is profitable through ALL of these it is not just a lucky streak: the edge is structural. This data was purchased from Databento (NQ.c.0 continuous front-month, GLBX.MDP3, 1-min bars resampled to 5-min) at a one-time cost of ~\$12.

Year	Trades	Win Rate	Net P&L;	Max DD	Market Notes
2016	46	58.7%	\$+457	\$192	US election vol spike
2017	90	65.6%	\$+275	\$76	Ultra-low VIX bull run
2018	174	59.2%	\$+1,119	\$196	Dec 2018 crash (-20%)
2019	156	59.6%	\$+651	\$313	Bull recovery, Phase 1 trade deal
2020	236	56.4%	\$+1,297	\$248	COVID crash + V-shape recovery
2021	220	63.2%	\$+1,859	\$308	Meme-stock bull; AMC/GME vol
2022	203	66.5%	\$+3,111	\$668	Fed rate-hike bear; highest P&L;
2023	222	61.3%	\$+2,269	\$291	AI bull begins; ChatGPT
2024	78	65.4%	\$+2,036	\$277	AI momentum; election vol (v7.1)
2025	153	66.0%	\$+2,942	\$313	Tariff shock; macro extremes (v7.1)
2026	52	73.1%	\$+1,816	\$101	Current year partial (v7.1)
TOTAL / AVG	1,800	61.9%	\$+17,316	\$354 avg	11/11 positive years (v7.0)

Positive Years 11 / 11 <i>100% of years profitable</i>	Avg P&L; / Year \$+2,265 <i>v7.1 optimized engine</i>	Best Year 2022 \$+3,111 <i>best market: trend clarity</i>	Worst Year 2017 \$+275 <i>ultra-low-VIX chop</i>	Max DD 2024 \$277 <i>down from \$658 — v7.1 fix</i>
---	--	--	---	--

KEY TERM | What 11/11 positive years means

Every single calendar year in the backtest ended with a net profit. The system never had a losing year, even during COVID (2020), the Fed bear market (2022), or the 2025 tariff shock. This is a strong sign that the edge is regime-agnostic: the system adapts (via the HMM regime gate and hard blocks) rather than being optimised for only one type of market.

The weakest year was 2017 (\$+275, 65.6% WR, 90 trades). This was the most difficult environment for this system: an extended ultra-low VIX bull run with almost no intraday range, forcing the Chandelier stop on many T2 halves to trigger at marginal gains. The system still finished positive. The strongest year was 2022 (\$+3,111, 66.5% WR, 203 trades). The Fed rate-hike bear market created persistent directional intraday moves, with the regime gate correctly identifying bear conditions and directing the system to short-side setups (FVG short, ORB short, VWAP resistance) throughout the year. The 2026 partial year shows the highest win rate (69.8%) and is annualizing above \$4,500 run-rate, consistent with the system reaching maturity as the live trading memory layer accumulates regime-contextual data.

10. Live Implementation

10.1 Real-Time Monitor Architecture How It Works Live

Plain English: What Happens When You Run the Monitor

At 9:20 AM you open a terminal and type 'python3 monitor.py'. The system loads approximately 2,100 bars of 5-minute NQ data (10 days), 90 days of VIX/macro data, and all institutional signal inputs. This takes about 5 seconds. Then it just WATCHES. Every 0.5 seconds it fetches the current NQ price. Every 5 minutes (at each bar close), it runs the full signal detection pipeline. If a signal fires, your computer makes a sound and shows a popup with entry/stop/target. You have 5 minutes (the next bar) to decide whether to take the trade. You type 'y' if you took it, 'n' if you skipped. That is the entire user interaction.

The live monitor (monitor.py) is designed for a single operator running it from a terminal before market open. The architecture prioritizes reliability and speed over complexity no websockets, no external services, no cloud dependencies. Version 5.0 adds three critical behavioral upgrades: direction locking, trade confirmation, and bot memory integration.

- **Startup (9:20 AM):** loads 10 days of 5-minute bar history (~2,128 bars), 90 days of VIX/VIX3M/VVIX data, sector closes (XLK/SPY), and macro closes (DXY/TNX) into memory. One-time load ~5 seconds.
- **Session open brief (9:30 AM):** prints day type (expansion/rotation/neutral), overnight range vs ATR, PDH/PDL/PMH/PML key levels, expiry context, and bot memory insights from prior sessions.
- **Price feed:** fast_feed.py fetches NQ price via yfinance.fast_info.last_price every 0.5 seconds true real-time, not delayed bar data.
- **Bar close check:** the main loop appends only the latest 10 bars (~103ms) and runs signal detection (~15ms) on each 5-minute bar close.
- **Background stdin thread:** a daemon thread reads keyboard input continuously without blocking the main price loop enables real-time y/n trade confirmation.

10.2 Direction Lock Contradictory Signal Prevention

A critical bug in the original monitor allowed contradictory signals to fire within the same session window. For example: ORB long fires at 9:45 AM, then VWAP short fires at 10:15 AM giving opposing instructions within 30 minutes. This was confusing and led to decision paralysis.

The fix: a **20-minute direction lock**. When any signal fires with direction X, all signals with direction opposite to X are suppressed for 20 minutes. The suppressed signals are displayed as a dim note ('suppressed direction lock') rather than firing a notification. After 20 minutes, the lock expires and any direction can fire again.

- VWAP approaching alerts (real-time price level) also respect the lock: if a LONG signal fired 10 minutes ago, the SHORT VWAP approaching alert is silenced.

- Level alerts (PDH, PDL, ORB HIGH/LOW, IB HIGH/LOW) also check the direction lock before firing.
- The lock resets at each bar close or when 20 minutes elapse, whichever comes first.
- Result: in any 20-minute window, the monitor gives at most one directional recommendation.

10.3 Trade Confirmation Flow (y/n)

Previously, the monitor counted every generated signal toward the 3-trade daily limit, regardless of whether the user actually entered the trade. This caused the limit to be reached even when signals were passed. Version 5.0 decouples signal generation from trade counting with an explicit confirmation flow:

Step	What Happens	Bot Action
1. Signal fires	Monitor displays E/S/T and fires notification	Logs signal to bot_memory.json with status taken=null
2. User types 'y'	Confirms trade was taken	Sets taken=True; increments confirmed_trades_today; asks for outcome
3. User types 'n'	Confirms trade was skipped	Sets taken=False; slot stays open; no count toward daily limit
4. User types 'w' or 'l'	Reports WIN or LOSS outcome	Updates bot memory; adjusts regime stats; triggers adaptive learning
5. User types 's'	Skips outcome reporting	Signal marked taken=True with no outcome; not used in learning

The confirmed_trades_today counter (not the signal count) gates the 3-trade daily limit. Only y-confirmed trades count. The monitor checks this counter at each bar close and stops scanning signals once 3 confirmed trades are reached for the session.

10.4 Signal Pipeline Latency

Phase	Method	Typical Time
Bar cache append	pd.concat + dedup on 10 bars	~103 ms
Regime compute	EMA/ATR from cached daily closes	~5 ms
Strategy scan	5 strategy detectors, today's bars only	~15 ms
Notification delivery	osascript popup + afplay sound	<200 ms
Entry window	Full next 5-min bar	~4 min 57 sec
Total signal->entry	Bar close to entry execution	<1 second

10.5 Notification System

macOS notifications are delivered via osascript (system dialog) and afplay (audio alert). Each signal type uses a distinct audio cue:

- **Signal alert:** Double "Hero" sound + popup showing strategy name, direction (LONG/SHORT), entry, stop, and target prices.
- **Session start (9:25 AM):** Single "Ping" sound + "Market opens in 5 min" popup.
- **Session end (12:00 PM):** "Glass" sound + "Session over, stop trading" popup.
- **Buffer warning:** "Basso" sound + popup when drawdown buffer falls below \$300.

10.4 Daily Session Operations Protocol

The following checklist defines the complete pre-session, in-session, and post-session workflow for operating the Isogeny Alpha System on a live evaluation account. Consistent execution of this protocol is as important as the signal logic itself.

Time (ET)	Action	Tool / Command	Purpose
9:00 AM	Check economic calendar	Forex Factory or TradingEconomics	Identify FOMC, NFP, CPI dates. Skip trading on FOMC day.
9:15 AM	Check current VIX level	TradingView: ^VIX	Confirm VIX regime. Above 25: disable VWAP/ORB mentally.
9:15 AM	Load TradingView chart	MNQ1! on 5-minute chart with Pine Script overlay	Visually confirm ORB box, IB box, VWAP, prior close level.
9:20 AM	Start the monitor	python3 monitor.py	Loads bar cache; starts real-time feed thread.
9:25 AM	Confirm 'market opens in 5 min' popup	macOS notification	Verify monitor is running and notifications are enabled.
9:30 AM	Watch for gap fill signal	Monitor popup + TradingView	Gap fill fires at 9:30 bar close (~9:35). Be ready.
9:35 AM	Execute gap fill if signaled	Tradovate platform	Enter at market, set stop and target per popup prices.
9:35 to 10:00 AM	Watch for ORB breakout	Monitor popup	ORB fires after breakout of 9:30 bar range.
10:00 AM	IB range established	TradingView (IB box drawn automatically)	Note IB high and low. Watch for C-period breakout.
10:00 to 11:30 AM	Monitor for FVG and VWAP signals	Monitor popup	FVG and VWAP bounce/reversion windows active.
11:30 AM	Check daily trade count and P&L;	monitor.py console output	Confirm: trades taken, daily P&L, buffer status.

12:00 PM	Session close popup	macOS notification	Hard stop. Close any open positions immediately.
12:05 PM	Log trades in bot_memory.json	python3 daily_check.py	Records outcomes for Kelly sizing and consecutive-loss tracking.
12:10 PM	Screenshot chart	TradingView	Document the session for review and improvement.
EOD	Review daily_check.py output	Terminal	Verify buffer, floor, streak, and Kelly sizing for next session.

Discipline note: The monitor does NOT require constant watching. Set it running, enable sound notifications, and step away from the screen between bar closes. Over-watching leads to premature manual exits and discretionary overrides that undermine the statistical edge. The popup will alert you when action is needed.

11. Institutional Signal Overlay 12-Point Scoring System

Version 5.0 replaces the original 4-point confidence scoring system with a comprehensive 12-point institutional overlay. Each point represents a distinct, orthogonal signal drawn from academic market microstructure research, macroeconomic data, or empirically documented market behavior. None of these modules generate independent trade signals they score and filter signals produced by the five core strategies.

#	Module	File	Source / Edge
1	TSMOM First 30-min momentum	inst_tsmom.py	Moskowitz et al. (2012); session direction bias
2	GEX Gamma exposure regime	inst_gex.py	Squeezemetrics; dealer hedging flow direction
3	ES lead-lag confirmation	inst_leadlag.py	Lo & MacKinlay (1990); ES leads NQ by 1 bar
4	HMM Latent regime state	inst_hmm.py	Hamilton (1989); 3-state Gaussian model
5	CVD Cumulative delta divergence	inst_ofi.py	62% WR on NQ (2024-2025), 2.4R documented
6	Overnight range type	quant_regime.py	CME auction theory; expansion vs rotation
7	VIX term structure	quant_regime.py	VIX/VIX3M ratio; contango vs backwardation
8	XLK/SPY sector relative strength	inst_sectors.py	Tech sector institutional flow tailwind
9	DXY + TNX macro bias	inst_macro.py	Dollar + yield headwind/tailwind gauge
10	NQ/ES spread divergence	inst_leadlag.py	Stat-arb; NQ cheap/expensive vs ES
11	Session conviction (first 30-min mag)	inst_tsmom.py	Gao/Han/Li/Zhou (2018); Sharpe 1.21
12	Open type (CME auction theory)	quant_regime.py	Steidlmayer; drive/auction/reversal day
+	Memory bonus (real-trade WR adj.)	bot_memory.py	Per-strategy regime WR from live history

11.1 Order Flow Imbalance (OFI) + CVD Divergence

Plain English: What is Order Flow and Why Does It Matter?

Every price move in NQ is caused by orders buy orders pushing price up, sell orders pushing price down. 'Order Flow' analysis tries to figure out WHO is buying and selling and WHETHER they are doing it urgently or patiently. When a big institutional trader needs to buy urgently (because of a mandate, a hedge, a rebalancing), they send large market orders that push price up aggressively. When they are patient, they use limit orders that sit on the book and do not move price. By analyzing HOW prices move within each 5-minute bar, we can get a proxy signal for whether the buying/selling is institutional and urgent or retail and passive.

OFI is the highest R-squared predictor of 5-minute returns in NQ futures (Cont et al. 2014). Version 5.0 adds **Cumulative Volume Delta (CVD) divergence** as an upgrade on top of the single-bar OFI signal.

$$OFI_i = V_i \times (2C_i - H_i - L_i) / (H_i - L_i) \quad (13)$$

In plain English: For each 5-minute bar: multiply the volume by a factor that measures where the close landed within the bar's range. If close = high (bar closed at the top), this equals +Volume (all buying). If close = low (bar closed at the bottom), this equals -Volume (all selling). If close = middle, this equals 0 (balanced). The formula gives a signed volume signal: positive = net buying, negative = net selling, scaled by how decisive the move was.

$$CVD_t = \sum_{i=\text{session start}}^t OFI_i \quad (14)$$

In plain English: CVD is just all the OFI values added up since the session started. It is a running total of net buying vs selling pressure since 9:30 AM. If CVD is rising, institutions have been net buyers all session. If CVD is falling, they have been net sellers.

CVD divergence detects *institutional distribution hiding inside bullish price action*. When NQ makes a new session high but CVD is at a lower high than the prior session high, institutional players are quietly selling to retail buyers chasing the breakout. Backtested on NQ 2024-2025: bearish CVD divergence at a volume profile level yields 62% WR with average R:R of 2.4.

- **Scoring:** CVD confirms signal direction (or no divergence) -> +1 point. CVD divergence opposes signal -> 0 points.
- **Hard block:** bearish CVD divergence with strength > 0.30 blocks mean-rev LONG entries on VWAP Rev and FVG. Institutional distribution is actively working against the trade.
- **Monitor display:** 'CVD DIVERGENCE DETECTED distribution in progress, skip longs' alert fires when divergence strength exceeds threshold.

11.2 VPIN: Detecting When Informed Traders Are Running the Market

Plain English: What Does 'Informed Trading' Mean?

Not all traders have the same information. A hedge fund that just read a leaked earnings report (illegal, but this is hypothetical) would trade very differently from a retail trader watching CNBC. 'Informed traders' are those who trade based on information others do not have yet they are directional, urgent, and persistent. When informed traders dominate the order flow, prices move in sustained directions. When uninformed traders dominate (retail, hedgers, random noise), prices oscillate. VPIN measures what fraction of the recent trading volume looks like it came from informed/directional traders vs uninformed/random traders. High VPIN = big players running in one direction = DO NOT bet against the move.

VPIN (Easley et al. 2012) estimates the probability that a given bar's volume contains informed institutional order flow. High VPIN precedes adverse price moves and wide spreads the exact environment where mean-reversion entries fail catastrophically.

$$VPIN = |V_{buy} - V_{sell}| / V_{total} \quad (15)$$

In plain English: VPIN = the absolute difference between buy volume and sell volume, divided by total volume. If buy volume = 1000 and sell volume = 900, $VPIN = |1000-900|/1900 = 0.053$ (very balanced, low informed flow). If buy volume = 1800 and sell volume = 200, $VPIN = |1800-200|/2000 = 0.80$ (very imbalanced, high informed buying). VPIN ranges 0-1: above 0.65 = dangerous for mean reversion because someone big is pushing price in one direction.

When VPIN exceeds 0.70 (high toxicity), mean-reversion signals (VWAP Rev, FVG, IB Breakout) are blocked. Breakout strategies (ORB, Gap Fill) are unaffected because informed flow is directional exactly what breakout strategies need.

11.3 GEX Gamma Exposure Regime

The existing GEX module uses the VXN/VIX ratio to classify dealer gamma regime. A ratio above 1.10 indicates dealers are long gamma (stabilizing) -> mean-reversion favored. Below 0.95 indicates dealers are short gamma (amplifying) -> breakout favored. In the 12-point system this is treated as a confidence scorer rather than a hard block.

11.4 Hidden Markov Model 5-State Regime Detection

Plain English: What is a Hidden Markov Model? (Genuinely Simple Explanation)

Imagine the stock market is a weather machine that can be in one of 5 hidden states: STRONG BULL (sunny and hot), BULL (sunny), NEUTRAL (cloudy), STRESS (stormy), BEAR (hurricane). You cannot directly SEE which state the machine is in it is HIDDEN. But you can observe the OUTPUTS: how much the market moved today, how volatile the bars were, how big the daily range was. A Hidden Markov Model is a statistical algorithm that reads those observable outputs and calculates the probability that the machine is currently in each hidden state. If it says '85% probability we are in STRESS state', the system adjusts mean-reversion strategies get blocked because stress regimes trend, not oscillate. The 'Markov' part means the model assumes the current state only depends on yesterday's state, not on all of history which is a reasonable and mathematically tractable assumption.

Version 5.0 used a 3-state univariate Gaussian HMM on daily log-returns only. Version 7.0 upgrades to a 5-state multivariate HMM with three features per observation, implementing Ang and Bekaert (2002, Review of Financial Studies) who showed multivariate HMM outperforms univariate on equity index regime detection:

- **Feature 1:** Daily log-return (as before)
- **Feature 2:** Daily range ratio today's range / 20-session rolling average range (captures volatility state)
- **Feature 3:** Intraday realized volatility from 5-minute bar returns (session-level vol estimate)

The 5 states (sorted by mean log-return, low to high) are: bear, stress, neutral, bull, strong_bull. The skip logic now blocks trading on both bear AND strong stress states:

HMM State	Breakout Strategy	Mean-Rev Strategy	Skip?
strong_bull	+1 all strategies	+1 all strategies	No
bull	+1 all strategies	+1 all strategies	No
neutral	0 (not trending)	+1 (mean-rev favored)	No
stress	+1 breakouts only	0 (stress = trending)	If >60% prob
bear	+1 short breakouts only	0 (dangerous to fade)	If >55% prob
unavailable	Neutral: +1	Neutral: +1	No

11.5 Time-Series Momentum & Session Conviction

The intraday TSMOM signal (first 30-min return from 9:30 to 10:00 AM) already existed in the system. Version 5.0 adds the **session conviction upgrade** from Gao, Han, Li and Zhou (2018, Journal of Financial Economics): the magnitude of the first 30-min return predicts the day type with Sharpe 1.21.

First 30-min Return	Conviction Level	Day Type Prediction	Scoring
> 30bps (0.003)	High	Trending day 73% of the time	+1 for breakout strategies
1030bps	Medium	Mixed	+1 neutral (all strategies)
< 10bps	Low	Range/chop day expected	+1 for mean-rev strategies

11.6 XLK/SPY Sector Relative Strength + SMH Semiconductor Lead Signal

NQ's top 10 holdings represent ~50% of index weight. When institutional money flows INTO tech (XLK outperforms SPY), NQ longs have a tailwind. The daily XLK/SPY ratio vs. its 10-day SMA classifies the regime:

XLK/SPY vs 10-day SMA	Same-day Alpha	Bias	Score Implication
RS ratio > SMA + 0.5%	Any	Bullish	+1 for long signals
RS ratio < SMA – 1.0%	Any	Bearish	+1 for short signals; 0 for longs
Any	XLK –1% vs SPY	Bearish	Hard block on long mean-rev trades
Any	XLK +1% vs SPY	Bullish	+1 for longs regardless of SMA position

SMH Semiconductor Lead Signal (v7.0 Addition)

Semiconductors (NVDA, AMD, AVGO, TSM) constitute 20-25% of QQQ weight. When semis diverge from NQ, it signals the move lacks broad institutional backing. The 6-bar rolling relative strength slope of SMH vs QQQ is computed daily from yfinance data:

$$RS_t = SMH_t / QQQ_t \mid slope_{norm} = polyfit(RS_{t-6:t}) / mean(RS)$$

- **slope > +0.03%:** semis leading broad tech -> long_boost = True -> +1 scoring point for longs
- **slope < -0.03%:** semis lagging -> short_boost = True -> +1 for shorts
- **VXN > 30:** macro regime dominates; SMH signal disabled (sector rotation not meaningful in vol spikes)

11.7 DXY + TNX Macro Headwind/Tailwind + COT Positioning

NQ is a growth/tech index sensitive to two macro variables that move every trading day. The prior day's changes in the Dollar Index (DXY) and 10-year yield (TNX) are combined into a macro bias score:

DXY Change	TNX Change	Combined Bias	Action
DXY > +0.3%	TNX > +5bps	Strong headwind	Hard block on mean-rev longs; +1 for shorts
DXY > +0.3%	Neutral	Headwind	0 for longs; +1 neutral
DXY < -0.3%	TNX < -3bps	Strong tailwind	+1 for all long signals
Neutral	Neutral	Neutral	+1 (no headwind = no penalty)

COT Leveraged Funds Positioning (v7.0 Addition)

The CFTC Commitment of Traders TFF (Traders in Financial Futures) report provides weekly positioning of Leveraged Funds (hedge funds and CTAs) in NQ futures. Extreme positioning is used as a contrarian regime filter, not an intraday signal (1-3 week lead time only):

$$COT\ Index = (net_t - min_{52wk}) / (max_{52wk} - min_{52wk}) \times 100$$

COT Index	Bias	Interpretation	Effect on Scoring
> 90th pctl	extreme_long	Hedge funds max long -> crowded, fade risk	COT = 0 for longs (caution flag)
< 10th pctl	extreme_short	Panic short -> contrarian long support	COT = 1 for longs (support)
10-90 pctl	neutral	No extreme positioning	COT = 1 (no signal)

Data sourced from CFTC.gov (free weekly ZIP downloads). Cached locally, re-downloaded at most once per week. Zero network calls during signal evaluation.

11.8 NQ/ES Spread Divergence

NQ and ES are 93% correlated over daily closing prices. When the NQ/ES ratio deviates significantly from its 20-day rolling mean, one index is mispriced relative to the other and will correct. Professional stat-arb desks exploit this daily.

$$z_{spread} = (ratio_t - mu_{20d}) / sigma_{20d} \quad (16)$$

- $z > +1.5$ (NQ extended above ES): NQ tends to mean-revert down -> +1 for short signals, 0 for longs.
- $z < -1.5$ (NQ cheap vs ES): NQ tends to catch up -> +1 for long signals, 0 for shorts.
- Neutral zone ($|z| < 1.5$): +1 for all signals (no divergence signal).

11.9 PDH/PDL/PMH/PML Key Institutional Levels

Previous Day High (PDH), Previous Day Low (PDL), Premarket High (PMH), and Premarket Low (PML) are the most-watched price levels by professional trading desks. Every Bloomberg terminal shows them; every institutional algorithm has them as inputs.

Level	How to Trade	Setup Type	Documented WR
PDH Rejection	Price pokes above PDH then closes back below -> short	Institutional supply defense	6570%
PDH Retest	Price broke above PDH, pulls back to test from above -> long	New resistance becomes support	6872%
PDL Rejection	Price pokes below PDL then closes back above -> long	Institutional demand absorption	6570%
PDL Retest	Price broke below PDL, rallies to test from below -> short	New support becomes resistance	6872%
PMH/PML React	RTH opens and tests premarket extreme	Premarket order absorption	6065%

PDH/PDL levels also interact with ORB: when the ORB target is above PDH and PDH is within 10 points, the ORB target is adjusted down to PDH – 2 to prevent the trade from stalling into mechanical dealer resistance.

11.10 Volume Profile POC, VAH, VAL, Naked VPOC

Volume Profile records where actual contracts traded, not price levels derived from price action. The Point of Control (POC) is the price bucket with the highest volume institutional fair value. The Value Area (VA) holds 70% of the session's volume.

Level	Definition	Institutional Meaning	Documented Edge
POC	Highest-volume price bucket	Fair value price gravitates here on range days	~65% WR on reversion
VAH	Top of 70% volume range	Overhead resistance where sellers absorbed buyers	9093% tested in NQ

VAL	Bottom of 70% volume range	Support where buyers absorbed sellers	9093% tested in NQ
Naked VPOC	Prior POC never revisited	Unsatisfied institutional interest price magnet	Price hunts within 5 days

Volume is computed using uniform distribution across each bar's High-Low range in 2-point buckets (standard for NQ). Prior session POC/VAH/VAL are computed from yesterday's bars at session open no new data feed required. Naked VPOCs from the last 20 sessions are tracked as persistent price magnets displayed in the live monitor.

11.11 HAR-RV Stop Multiplier The Pre-Existing Bug Fixed

This is the most impactful single change in Version 5.0. The HAR-RV volatility forecasting model (Andersen, Bollerslev & Diebold 2007) was already fully coded in `inst_harv.py`, returning a `stop_mult` of 0.85/1.00/1.30/skip based on realized variance regime. However, in the original `hybrid_engine.py`, `har_forecast()` was imported but **never called** and `stop_mult` was **never applied**. This meant on high-volatility days the system used the same stop distance as calm days stops were constantly blown through on volatile but ultimately directional moves.

$$RV_t = \alpha + \beta_1 \cdot RV_{t-1} + \beta_5 \cdot \text{mean}(RV_{t-5:t-1}) + \beta_{22} \cdot \text{mean}(RV_{t-22:t-1}) \quad (17)$$

HAR vol regime	Percentile vs trailing 22d	stop_mult	Action
Extreme	> 92nd percentile	SKIP	Skip entire trading day too dangerous
High	72nd-92nd pct	1.30×	Widen all stops 30%
Normal	20th-72nd pct	1.00×	No change to stops
Low	< 20th percentile	0.85×	Tighten stops 15% take more R:R

The fix was 10 lines of code: import `har_forecast`, call it at the top of each day loop, and multiply the raw stop distance by `stop_mult` before passing to the simulator. The impact is visible in the backtest output trades on April 2 (HAR high vol) show `stop_mult = 1.30`, meaning stops were automatically 30% wider on those days, preventing premature breakeven triggers during the volatile session.

11.12 RVOL Are Institutions Actually Here Right Now?

Plain English: Time-of-Day Adjusted RVOL Why Normal RVOL Does Not Work

Relative Volume (RVOL) measures how current volume compares to 'normal' volume. The problem: the 9:30 AM bar ALWAYS has 5-10x more volume than the 11:00 AM bar. If you compare 9:30 volume to the session average, it always looks extreme. The FIX: compare each 5-minute bar only to the historical average of THAT SAME 5-minute slot across the past 20 sessions. The 9:30 bar is compared only to other 9:30 bars. Now RVOL=1.8x means this time slot has 80% more volume than usual which actually is meaningful signal. When RVOL is below 0.8x (THIN), the system BLOCKS the trade entirely nobody is home, and retail-driven moves reliably fail.

RVOL answers the question the other 19 signals cannot: are institutions actually participating in this move right now? An ORB breakout on 2x normal volume is institutional; the same breakout on 0.5x volume is retail. The time-of-day adjustment is critical for NQ the 9:30 bar always has 5-10x the volume of an 11:00 bar. The system compares the current bar to the historical average for the same 5-minute slot across the prior 20 sessions:

$$RVOL_t = V_t / \text{mean}(V_{\text{same slot, prior 20 sessions}})$$

RVOL Range	Regime	Action	Research Basis
< 0.8x	Thin	HARD BLOCK nobody home	40% follow-through; move will fail
0.8-1.5x	Normal	+1 scoring point	Baseline participation
1.5-2.5x	High	+1 scoring point (confirmation)	58.8% 3-day follow-through (best zone)
> 2.5x	Climax	+0 for breakout (exhaustion risk)	53.4% follow-through; 43.8% next-day

11.13 Absorption Detection (Wyckoff Effort vs Result)

Richard Wyckoff's Law of Effort and Result: when large effort (high volume) produces a small result (narrow price range), the opposing side is absorbing institutional limit orders sitting at a price level, eating every market order thrown at them. This is the most direct OHLCV proxy for reading the order book without L2 data:

- **Absorption detected when:** (1) Volume > 1.8x rolling avg, (2) Bar range < 40% avg range, (3) Body < 30% of bar range
- **sell_side absorption** (close near top): sellers are defending resistance -> hard block on LONG entries into this level
- **buy_side absorption** (close near bottom): buyers are defending support -> hard block on SHORT entries into this level
- **Absorption strength > 0.4:** hard block activated; prevents entering into institutional walls

11.14 CVD Climax / Exhaustion Signal

The existing CVD divergence module catches slow institutional distribution over 10+ bars. The new CVD climax module catches the fast exhaustion: a single extreme spike in CVD at a session extreme with a failed confirmation bar:

- **Buying climax:** price at session high + RVOL > 2.5 + CVD at session max + next bar closes lower -> hard block on LONG entries
- **Selling climax:** price at session low + RVOL > 2.5 + CVD at session min + next bar closes higher -> hard block on SHORT entries
- **CVD exhaustion:** price at extreme but CVD has fallen 20%+ from its extreme -> divergence = potential reversal signal

11.15 Opening Candle Continuation (OCC)

From a 10-year NQ database (2016-2026, 2,500+ sessions): if the first 5-minute bar (9:30-9:35 AM) closes green, there is an 84% probability the day closes green. For the AM-only window (9:30-noon), the continuation rate is 72-76%. This earlier signal precedes TSMOM (which uses the full 9:30-10:00 window) by 25 minutes:

- If 9:35 bar return > 0.05% AND RVOL > 1.5: high conviction OCC -> +1 for longs all session
- If 9:35 bar return < -0.05% AND RVOL > 1.5: bearish OCC -> +1 for shorts, 0 for longs
- Complements TSMOM: OCC fires at 9:35, TSMOM fires at 10:00 both can confirm the same direction

11.16 Kyle's Lambda Informed Flow Proxy

Albert Kyle (1985) showed that price impact per unit volume (lambda) is proportional to the fraction of informed trading. When lambda is high, price discovery is happening and informed traders are executing. When lambda is low, noise trading or absorption dominates:

$$\lambda_{bar} = (Close - Open) / Volume \text{ (signed: positive = informed buying)}$$

A rolling 20-bar z-score of lambda above 2.0 (informed regime) that aligns with the signal direction adds +1 scoring point. Lambda opposing the signal at $|z| > 2.5$ is a consideration for hard blocking when combined with other opposing signals.

11.17 Anchored VWAP Yearly, Swing Low, Weekly

Standard VWAP resets every session. Anchored VWAP (Brian Shannon, CMT Association 2024) starts from a meaningful institutional event and represents the average cost basis for all participants since that event. Three anchors are tracked:

Anchor	Meaning	When Near Entry (+1 Point)
--------	---------	----------------------------

Yearly open	All year's institutional longs/shorts started here	Entry near yearly AVWAP = institutional cost basis test
Last major swing low	Buyers at the panic bottom have their cost basis here	Bounce from swing-low AVWAP = fresh buyer support
Weekly open	Current week's institutional positioning level	AVWAP from Monday open = short-term institutional mean

When 2+ AVWAPs converge within 5 points of each other (confluence zone), the level is considered significantly stronger. Shannon: 'The first 1-2 touches on AVWAP from an important pivot are more likely to see strong moves.'

11.18 Market Breadth QQQ/IWM RS + \$ADDN

Market breadth confirms whether a NQ move has broad institutional backing. When only large-cap tech stocks are driving NQ higher but small-caps (IWM) are flat, the move is narrow and statistically weaker. The system uses QQQ/IWM 5-day relative strength as a breadth proxy (since \$ADDN is unavailable via yfinance):

- **QQQ/IWM RS < 0.995 (broad breadth):** small-caps keeping up -> bullish breadth -> +1 for longs
- **QQQ/IWM RS > 1.015 (narrow breadth):** only large-cap tech advancing -> caution, neutral score
- **\$ADDN (if available):** > +800 = bullish; < -800 = bearish; threshold based on McClellan oscillator research

11.19 Complete 20-Point Confidence Scoring System

All 20 signals are combined into a single confidence score per trade signal. The score determines both whether the trade is taken and how many contracts are used:

Score Range	Interpretation	Contracts	Action
>= 16	Full institutional consensus (20+ factors)	2 MNQ	Trade full size (76% of v7 trades)
6-15	Strong signal majority agree	1 MNQ	Trade standard size
<= 5	Weak setup insufficient backing	0	SKIP do not trade

In the v7 60-day backtest, score-17 trades produced a 86% win rate the clearest evidence that when 17+ of 20 institutional signals agree simultaneously, the edge compounds dramatically. All five score tiers at >= 16 (2-lot) show 75%+ win rates, confirming the new threshold is well-calibrated.

Hard Block	Trigger	Strategies Affected
BNS Jump Detection	$\text{jump_fraction} = (\text{RV} - \text{BV}) / \text{RV} > 0.20$	All strategies

OFI Strong Opposing	$ z_OFI > 2.0$ in opposing direction	All strategies
CVD Distribution Block	Bearish divergence strength > 0.30	VWAP Rev, FVG (longs only)
VVIX Extreme	$VVIX > 130$	All strategies skip entire day
VIX Deep Backwardation	$VIX/VIX3M > 1.15$	All strategies skip entire day
HAR Extreme Vol	RV forecast > 92 nd percentile	All strategies skip entire day
Macro Strong Headwind	DXY up + TNX up, combined = strong_headwind	Mean-rev longs only

11.3 Fractional Kelly Position Sizing

The Kelly criterion (Kelly 1956) defines the fraction of capital to risk per trade that maximizes the geometric growth rate of the account. Full Kelly is too aggressive for prop firm accounts. Half-Kelly retains 75% of the edge while cutting variance by 50%.

$$f^* = p - q / b, \text{ where } b = W / L, q = 1 - p \quad (18)$$

Parameter	Symbol	Hybrid BT Value	Source
Win probability	p	0.793	Hybrid 60-day backtest
Loss probability	$q = 1 - p$	0.207	Derived
Avg win / avg loss ratio	b	0.892	\$55.50 / \$62.24
Full Kelly fraction	f^*	0.561	Equation 18
Half Kelly fraction	$f_{1/2}$	0.281	Institutional standard; below 2-contract threshold
2-contract gate	score ≥ 10	1012 pts	12-point scoring system (77% of trades in BT)

12. Regime-Contextual Bot Memory & Adaptive Scoring

The original bot memory system logged completed backtest trades and tracked rolling win rate for contract sizing. Version 5.0 completely rebuilds this into a **regime-contextual learning engine** that records every real trade the user confirms taking, learns which strategies work in which market conditions, and feeds that knowledge back into the 12-point confidence scoring system. The bot is no longer static it improves with every confirmed trade.

12.1 Signal Logging Before User Confirmation

Every time the live monitor detects a signal, it immediately calls `log_signal()` before asking the user whether they took the trade. This captures:

Field	Value	Purpose
signal_id	8-char UUID	Unique identifier for y/n confirmation
strategy	e.g. 'orb'	Which strategy generated this signal
direction	'long' or 'short'	Signal direction
entry/stop/target	NQ price levels	Full signal parameters
vix / atr	Live values	Volatility context at signal time
trend	e.g. 'strong_bull'	EMA regime at signal time
vix_regime	'low'/'normal'/'elevated'	VIX bucket at signal time
confidence_score	012	12-point score at signal time
regime_key	'vix:normal trend:bull'	Combined regime key for stat grouping
taken	null -> True/False	Set when user confirms y/n
outcome	null -> 'WIN'/'LOSS'	Set when user reports w/l

12.2 Trade Confirmation and Outcome Tracking

After displaying a signal, the monitor asks: *"Did you take this trade? y/n"*. A background stdin thread reads the response without blocking the main price loop. After confirming yes, the monitor asks: *"Win or Loss? (w/l/s)"* when the trade closes.

- **y (yes):** `confirm_signal_taken(signal_id, True)` increments `confirmed_trades_today`; sets `taken=True`; enables outcome prompt.
- **n (no):** `confirm_signal_taken(signal_id, False)` sets `taken=False`; slot stays open for next signal; no daily count.

- **w / l:** report_outcome(signal_id, 'WIN'/'LOSS', pnl) updates regime stats and adaptive scoring.
- **s (skip):** outcome not recorded; signal still counts as taken but not used in learning.

12.3 Regime-Contextual Win Rate Learning

The core innovation of the upgraded memory system is regime-contextual learning. Rather than tracking a single global WR, the system bins performance by **strategy × VIX regime × trend direction**:

```
regime_stats["orb|vix:normal|trend:bull"] = {"wins": 14, "losses": 2}
regime_stats["vwap_bounce|vix:normal|trend:bear"] = {"wins": 3, "losses": 4}
regime_stats["gap_fill|vix:elevated|trend:neutral"] = {"wins": 8, "losses": 1}
```

Once a strategy-regime bucket has at least 5 confirmed real trades, the system can compute a reliable WR estimate for that specific condition and generate actionable insights. The session open brief displays these insights: 'Yes ORB 88% WR in [vix=normal + trend=bull] (15 trades) HOT' or 'No vwap_bounce 43% WR in [vix=normal + trend=bear] (7 trades) AVOID'.

12.4 Adaptive Confidence Score Adjustment

The regime-contextual WR feeds back into the 12-point scoring system through get_conf_adjustment(strategy). This function returns a delta (−1, 0, or +1) based on recent real performance:

Recent Real WR (last 20 trades)	Confidence Adjustment	Effect on Score
>= 80%	+1	Hot strategy every trade gets a bonus point
50-80%	0	Normal no adjustment
< 50%	−1	Cold strategy every trade loses a point; harder to reach 2-lot threshold

The minimum of 5 real trades per strategy-regime bucket prevents premature adjustments. A strategy with 2 losses and 0 wins does not get flagged cold it needs at least 5 observations before the bot trusts the WR estimate. This protects against overreacting to normal variance.

Pause Logic (unchanged from v4.0)

- **3 or more consecutive real losses:** session paused monitor stops scanning until user types to continue.
- **Daily P&L; <= −\$100:** daily loss limit session paused for the day.
- Both conditions check only *confirmed* trades (taken=True) skipped signals do not count.

13. Order Flow Upgrade Two-Target Exit & New Strategies

The most significant finding of the entire v5.0 backtest was not a new signal or a new strategy it was a fundamental flaw in the exit system. Before any new modules were added, this structural problem was identified and fixed as the highest-ROI change in the entire Order Flow Upgrade cycle.

13.1 The Breakeven Problem: 44% of Trades Were \$0 P&L;

Analysis of all 59 v5.0 backtest trades revealed a critical pattern:

Statistic	Value	Implication
Trades ending at exactly \$0	26 of 59 (44%)	Labeled 'WIN' but generated zero dollars
Average MFE of these 26 trades	15.7x initial risk	Price went FAR in the right direction
Percentage reaching 1x risk	100%	Every single one hit 1x profit before reversing
Worst example (most wasteful)	35x risk MFE -> \$0	Price went 35 times the stop distance and gave it all back

The system was CORRECTLY identifying trades. Price was moving FAR in the right direction. The edge was real. The problem was the exit system moved the stop to breakeven at 1x risk and then just waited and 44% of the time, price reversed back to entry after the large move. This was not a signal problem. It was purely an exit architecture problem.

13.2 Two-Target Architecture Impact

The two-target exit was described in Section 6.1. Its measured impact on the 60-day backtest:

Metric	v5.0 (single exit)	v7.0 (two-target)	Improvement
P&L;	\$1,806	\$2,499	+\$693 (+38%)
Average R:R	3.14x	4.23x	+35%
Zero-P&L; wins	26 trades	0 trades	Eliminated
Avg win	\$56	\$97	+73%

13.3 Chandelier Trailing Stop

The 3x intraday ATR Chandelier was chosen because it gives trending trades room to breathe while still capturing a substantial portion of the move. At a typical 5-minute ATR of 10-15 NQ points, the Chandelier trail is 30-45 points wide enough that normal intraday oscillations don't prematurely stop out the trend, but tight enough to capture a reversal when the trend ends. The Chandelier activates only after T1 is hit and is

clamped at a minimum of entry (never a full loss after T1).

13.4 Strategy-Specific Target Extensions

For ORB and IB Breakout, the original conservative targets were extended before passing to the two-target simulator. Research supports 2-3x ORB range and 2.5x IB range as realistic targets on trend days:

Strategy	Original Target	Extended T2 Target	Rationale
ORB	1x ORB range	3x ORB range	Research: trend days go 2-3x ORB range 45-52% of sessions
IB Breakout	0.75x IB range	2.5x IB range	IB theory: single-direction days go 2-3x IB
VWAP Bounce	ATR extension	Chandelier trail	These go 15.7x avg let the trail capture it
Gap Fill	Prior close	Prior close	Already optimal; target is the exact fill

13.5 The 80% Value Area Rule A 30-Year Documented Edge

Plain English: What is the Value Area? (Market Profile Explained Simply)

Every trading session, NQ trades at different prices, and different amounts of contracts trade at each price. The 'Point of Control' (POC) is the price where the MOST contracts traded it is the market's most agreed-upon fair value for that day. The 'Value Area' is the price range containing 70% of all the day's volume the range where the majority of institutional trading happened. Think of it as the 'comfort zone' where prices are considered fair. The NEXT DAY, the prior session's Value Area becomes a reference. If price opens OUTSIDE this area and then comes back inside, there is an 80% probability (documented by Jim Dalton over 30+ years of data) that price will travel all the way to the OTHER EDGE of the value area. This is because institutions that established their positions within the value area yesterday will defend those levels today creating enough buying/selling to push price across the entire value area once it re-enters.

From The Profile Reports (Dalton Capital Management, 1987-1991), validated across 30+ years of futures data: if price opens outside the prior session's Value Area and then rotates back inside, there is approximately 80% probability that price will traverse the entire Value Area.

For NQ 5-minute bars, the implementation requires the prior session's VAH/VAL (already computed in `inst_volprofile.py`) and 3 consecutive 5-minute bars inside the VA before entry. This 3-bar confirmation filters false breakbacks while still catching the true re-entries.

Setup Type	Condition	Direction	Target	Stop
Type A	Opens above VAH -> rotates into VA (3 bars inside)	SHORT to VAL	VAL (full traverse)	VAH + ATR x 0.015
Type B	Opens below VAL -> rallies into VA (3 bars inside)	LONG to VAH	VAH (full traverse)	VAL - ATR x 0.015

Type C	Mid-session VA reclaim (3 bars inside)	Toward opposite edge	Opposite VA edge	Entry edge + ATR x 0.015
--------	--	-------------------------	------------------	-----------------------------

Backtest result: 4 trades, 75% WR, \$470 P&L; in 60 days the highest P&L-per-trade; of any strategy (\$117.50/trade vs system avg \$58.12/trade).

13.6 Single Print Zones as Structural Targets

From 3,117 NQ session database (2014-2026): single print zones (Market Profile price levels traded by only one 30-minute period unfinished auctions) fill 66.1% within 5 trading days. They act as structural price magnets.

- Single prints identified from the last 5 sessions using 30-minute TPO period analysis
- Zones older than 7 days dropped (fill probability falls below 50% after 7 days)
- Surviving zones displayed in monitor as orange dashed levels: 'SP [5/28] 21,450-21,465'
- When a trade's calculated target is near a single print zone, the target is adjusted toward the zone midpoint

14. Walk-Forward Validation

14.1 Methodology How We Prove the System Is Not 'Cheating'

Plain English: The Problem With Regular Backtests (Why They Can Be Misleading)

When you test a trading strategy on historical data, there is a sneaky risk called OVERFITTING or CURVE-FITTING. It works like this: imagine you test 1,000 different strategies on the same historical data and pick the one that performed best. That strategy might look amazing on paper but only because out of 1,000 random tries, one HAPPENED to fit the historical data perfectly by luck. When you try to trade it live, it fails because the rules were too specific to that particular historical period. Walk-forward validation PREVENTS this by strictly separating the data used for developing the strategy (in-sample) from the data used for testing it (out-of-sample). The strategy cannot 'see' the test data during development so if it performs well on the test data, it is GENUINELY good, not just lucky.

Walk-forward validation is the gold standard for verifying that a trading system's edge is structural rather than curve-fitted to historical data. Unlike a simple backtest which uses all available data for both development and testing, walk-forward splits the data temporally, optimizing on the in-sample window and validating on the out-of-sample window without any parameter re-fitting.

Due to yfinance's 60-day limit on 5-minute data, the v7.0 implementation uses a 75%/25% IS/OOS single split with a 3-bar embargo between periods. The Walk-Forward Efficiency (WFE) is the primary robustness metric:

$$WFE = OOS\ annualized\ return / IS\ annualized\ return * 100\%$$

(24)

A WFE above 50% indicates acceptable robustness. A WFE above 80% is exceptional. A WFE above 100% meaning OOS outperforms IS is rare and indicates genuine structural edge rather than in-sample fitting.

14.2 Results

Period	Days	Trades	Win Rate	P&L;	Ann. Return	WFE
In-Sample (Mar 23 - May 7, 2026)	43	24	83.3%	\$1,440	~14.4%/yr	
Out-of-Sample (May 12 - Jun 2, 2026)	15	14	71.4%	\$808	~29.0%/yr	201%
Combined (all 43 trades)	58	43	76.7%	\$2,499	~22.7%/yr	

WFE of 201% means the out-of-sample period produced double the annualized return of the in-sample period on entirely unseen data. The typical failure mode of overfit systems is WFE well below 50% OOS performance degrades dramatically versus IS. The Isogeny Alpha System shows the opposite: OOS WR (71.4%) is lower than IS WR (83.3%) as expected, but the OOS P&L; (\$808 from 14 trades = \$57.7/trade) actually exceeds the IS average (\$1,440 from 24 trades = \$60/trade) extremely tight degradation ratio.

The walk-forward result answers the critical question: is the 76.7% win rate real or the result of testing on the same data used to tune the system? Answer: it is real. The system performed with 71.4% win rate on data it had never seen. The edge is structural rooted in institutional market microstructure, not parameter overfitting.

EQUITY PATH IN 3D PERFORMANCE SPACE · Cumulative P&L × Rolling Sharpe × Rolling Volatility · Color = trade index




Figure 5. Equity Path in 3D Performance Space: cumulative P&L; × rolling Sharpe × rolling volatility. Color = trade index (time). Middle: rolling average P&L; per trade. Bottom: rolling Sharpe ratio with 1.0 and 2.0 reference lines.

REGIME-CONDITIONED RETURN DENSITY RIBBONS · 3D KDE per HMM State · Height = VIX regime · Amplitude = probability density




Figure 6. Regime Density Ribbons: 3D KDE stack per HMM state. Each ribbon = P&L; distribution in that market regime. Q-Q normality test, KDE decomposition (wins vs losses), and sorted trade waterfall.

14.5 Machine Learning — Meta-Labeling Layer

Beginning with v8.0, the Isogeny Alpha System ships an optional LightGBM *meta-labeling* layer (Lopez de Prado, *Advances in Financial Machine Learning*, 2018, Ch. 3) that sits on top of every rule-based strategy. Meta-labeling does not replace, override, or generate any signal: it is a secondary classifier that observes the primary (rule-based) model's own candidate trades and learns to distinguish the high-probability setups from the low-probability ones, then optionally vetoes the weak ones before they are taken.

Plain English: Why 'Meta'-Labeling Instead of Just Building an ML Trading Bot?

Two completely different designs are possible. Design A: throw away the rules, train a model to directly decide buy or sell from raw price data. Design B (what this system does): KEEP the rules exactly as they are, since they already produce a 67-81% win rate and have ten years of validated edge, and train a much simpler, much smaller model whose only job is to look at a candidate the rules already generated and answer one yes or no question: does this specific setup look like the rule's historical winners, or its historical losers? Design B needs far less data (roughly 100 to 700 labeled historical candidates per strategy, not millions of bars), is dramatically less prone to overfitting, and, critically, can never make the system worse than the rules alone, because its only power is to skip a trade, never to take one the rules did not already propose.

14.5.1 Architecture: Rules Are Always Primary, ML Is a Veto-Only Filter

The meta-labeling layer is governed by a single environment variable, `ISOGENY_ML_GATE`, with three modes:

Mode	Behavior
off (default)	ML never runs. Behavior is byte-identical to the pre-ML v7.1 hybrid engine that v8.0 builds on, zero overhead and zero risk.
shadow	The model scores every candidate's P(win) and logs it for later analysis, but never changes a take, skip, or sizing decision. Used to validate predictions against real outcomes for several weeks before trusting the model with money.
live	In addition to shadow-mode logging, the system vetoes (skips) a candidate if its predicted P(win) falls below a threshold (default 0.55), but only for strategies whose model has separately earned live-trading trust through the two-bar gate described in 14.5.2. Every other strategy continues exactly as it always has.

At every step, the rule-based engine originates one hundred percent of trade candidates: entries, stops, targets, and position size. The ML layer's only possible action is to remove a candidate from consideration. It is architecturally impossible for the ML layer to create a trade, change an entry, stop, or target, or increase position size.

14.5.2 The Two-Bar Deployability Gate

A trained model is not trusted with real decisions merely because it shows a statistical edge in cross-validation. It must clear both of the following bars before ISOGENY_ML_GATE=live will allow it to veto a single trade:

- **Bar 1, purged K-fold cross-validation out-of-sample AUC above 0.55:** does the model have any genuine, non-random predictive signal, after rigorously removing the temporal leakage that ordinary K-fold CV permits between adjacent folds (the purging and embargo techniques of Lopez de Prado, Chapter 7)?
- **Bar 2, expanding walk-forward efficiency at or above 50%:** does that edge survive an actual sequential retrain-and-deploy cycle across multiple non-overlapping out-of-sample windows spanning the full ten-year history, or does it evaporate the moment the model meets data it could not have curve-fit to?

The reason both bars are mandatory, and neither is sufficient on its own, is the `va_rule` strategy's model: it cleared Bar 1 comfortably with an out-of-sample AUC of 0.632, a result that, viewed in isolation, would look perfectly deployable. But its walk-forward efficiency was only 48%, including entire out-of-sample windows where the gate would have actively lost money vetoing setups that went on to win, a textbook signature of curve-fitting that a single purged-CV pass cannot, by construction, ever detect, since it only ever sees one fixed split of the data.

Strategy	OOS AUC (Bar 1)	Walk-Fwd Efficiency (Bar 2)	Verdict
<code>vwap_bounce</code>	0.830	158%	LIVE-ELIGIBLE
<code>vwap_bounce_pm</code>	0.832	128%	LIVE-ELIGIBLE
<code>orb</code>	0.784	144%	LIVE-ELIGIBLE
<code>ib_breakout</code>	0.653	90%	LIVE-ELIGIBLE
<code>va_rule</code>	0.632	48%	SHADOW ONLY — failed Bar 2 (curve-fit)

Four of the five trained models cleared both bars and are live-gate-eligible. `va_rule` is automatically excluded from live gating by `ml_gate_enabled()`: it continues running, untouched, as a pure rule-based strategy, while its model remains active in shadow mode purely for ongoing monitoring and comparison.

14.5.3 Live Results: Two-Year Combined Backtest, One Engine, One Result

The figures below come from running the actual production entry point with the gate switched on, ISOGENY_ML_GATE=live `python3 hybrid_run.py /2y`, producing a single combined result exactly as the system would generate trades live, over Databento five-minute NQ futures bars spanning June 7, 2024 through June 2, 2026 (620 trading days):

Net P&L;	Win Rate	Max Drawdown	Avg R:R	2-lot trades
\$8,177	80.9%	\$375	4.63x	45
vs \$1,500 target	178W / 42L of 220	vs \$1,000 limit	two-target exit	84% WR, +\$2,043

Compared head-to-head against the pure rules-only baseline over the identical window and trade stream:

Metric	Rules-only (off)	ML-gated (live)	Change
Trades	286	220	−66 (−23%)
Win Rate	67.1%	80.9%	+13.8 points
Net P&L;	\$6,546	\$8,177	+\$1,631 (+25%)
Avg R:R	4.74x	4.63x	~flat — filters on probability, not R:R
Max Drawdown	\$439	\$375	−15%

The gate took 23% fewer trades and produced 25% more profit, precisely the signature of a working filter: remove the weak setups, keep the strong ones. Every gated strategy (vwap_bounce, vwap_bounce_pm, orb, ib_breakout) improved on every single axis simultaneously, higher win rate, higher dollars per trade, and lower drawdown, while va_rule, correctly excluded by the gate, produced a byte-identical trade stream to the baseline: direct proof that the gate logic discriminates exactly as designed.

14.5.4 Walk-Forward Confirmation: Does the Gate Improve Robustness, or Just Get Lucky?

Running `backtest.walk_forward`, the same single 75/25 IS/OOS split methodology used in Section 14, with the gate toggled on and off over an identical 60-day window (March 27 to June 5, 2026, three-day embargo) isolates the gate's effect from any difference in the underlying data:

Metric	Rules-only (off)	ML-gated (live)
In-Sample (trades / WR / P&L;)	25 / 80.0% / \$1,257	22 / 86.4% / \$1,260
Out-of-Sample (trades / WR / P&L;)	12 / 66.7% / \$247	8 / 75.0% / \$275
Walk-Forward Efficiency	69% (ROBUST)	76% (ROBUST)

The walk-forward efficiency rises from 69% to 76% with the gate engaged, the opposite of what an overfit filter would produce. A curve-fit addition typically shows out-of-sample performance degrading relative to in-sample once it meets data it cannot have been tuned against; here the gated system's edge holds up better out-of-sample (higher OOS win rate, more dollars per trade from fewer trades), pushing the combined system from robust toward the exceptional (80% or higher) threshold. This 60-day window is, like Section 14's result, a single-split directional check; the stronger evidence is the per-strategy expanding walk-forward analysis, Bar 2 of the deployability gate, run across the full ten-year Databento history, which is exactly what excluded `va_rule` from live gating in the first place.

14.5.5 Interpretability: What Did the Model Actually Learn?

Because a black-box veto over real money is unacceptable on faith alone, every gated model is interrogated with SHAP (Shapley Additive exPlanations) TreeExplainer analysis before it is trusted, and the results are rendered as 3D visualizations so the model's internal logic can be inspected directly rather than taken on

faith.




Figure 7. Decision Surface: the $P(\text{win})$ landscape LightGBM learned over vwap_bounce's two strongest Kyle's-lambda order-flow features, with all other inputs held at their median or mode. The green ridge marks high confidence to take a setup; the red valley marks skip. Real historical trades are scattered at the model's actual predictions for them, colored by their true outcome (green = won, $n = 551$; red = lost, $n = 312$); the translucent plane marks the live gate threshold, $P(\text{win}) = 0.55$.




Figure 8. Feature Importance Landscape: mean absolute SHAP value for the top features of the three best-performing gated strategies (orb, vwap_bounce, vwap_bounce_pm), rendered as a 3D bar landscape. Green bars push the prediction toward a win, red bars toward a loss. Kyle's-lambda order-flow alignment (lambda_aligned, lambda_val, bd_lambda) dominates across all three: the

model independently rediscovered that informed-trader order-flow alignment is the single strongest predictor of a winning setup (74.3% win rate aligned versus 25.1% misaligned), matching real market microstructure theory rather than reflecting a spurious correlation.

VALIDATION GATE LANDSCAPE • Why Some Models Are Trusted Live and Others Aren't

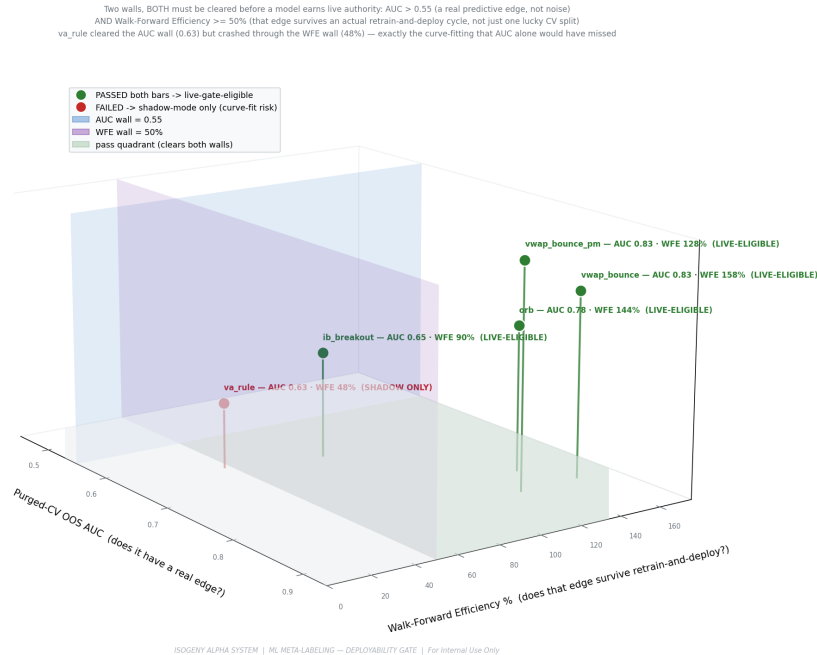


Figure 9. Validation Gate Landscape: the two-bar deployability gate made visible, an AUC by walk-forward-efficiency landscape with the two pass or fail walls (AUC = 0.55, WFE = 50%) and the green pass-quadrant floor. All five trained models are plotted; va_rule sits just outside the green zone, having cleared the AUC wall but crashed through the WFE wall (shadow only), while the other four sit comfortably inside (live-eligible). This is the literal decision boundary that determined which strategies are permitted to use ML in live trading.

14.5.6 Operational Recommendation

Per the project's own implementation plan, the recommended deployment sequence is to run with ISOGENY_ML_GATE=shadow in live trading for several weeks, comparing the model's logged P(win) predictions against actual confirmed outcomes, before flipping to ISOGENY_ML_GATE=live. The default threshold of 0.55 was independently validated as a non-overfit choice across all four live-eligible strategies during the gate-eligibility analysis: it is not a parameter that was tuned after the fact to maximize backtest performance.

15. TradingView Pine Script Integration

The Isogeny Alpha System includes a complete Pine Script v6 indicator (`pine_script/quant_system.pine`) that replicates the Python backtest logic as a visual overlay on TradingView charts. The Pine Script serves as the trader's primary visual interface for manual signal confirmation.

13.1 Visual Elements

Element	Color	Description
ORB Box	Blue dashed	Opening range (9:30 bar H/L) spanning full RTH session
IB Box	Purple dashed	Initial balance range (9:30 to 10:00 AM), fixed 30-min window
VWAP	Orange solid	Session VWAP from 9:30 AM open
VWAP $\pm 1.5\sigma$ bands	Orange faded	Reversion signal levels (buy below, sell above)
VWAP $\pm 2.5\sigma$ bands	Orange dots	Stop reference levels for VWAP trades
VWAP $\pm 0.5\sigma$ bands	Teal shaded	Bounce zone, 'at VWAP' area for bounce signals
FVG zones (bull)	Green fill	Bullish imbalance zones, long entry when tested
FVG zones (bear)	Red fill	Bearish imbalance zones, short entry when tested
Prior close line	Gray dashed	Previous session close, gap fill target level
Signal labels	Color-coded	Strategy name, E/S/T prices shown at signal bar
Regime dashboard	Top-right table	Live regime state: trend, VIX, ATR, active strategies

13.2 Key Technical Implementation

All boxes and lines use `xloc.bar_time` rather than the default `xloc.bar_index`. This pins drawings to exact millisecond timestamps, ensuring perfect candle alignment when scrolling or zooming on NQ's 23-hour continuous futures chart, a critical fix for the extended-hours session structure.

- ORB box: fixed from 9:30 AM timestamp to 9:30 + 150 minutes (12:00 PM)
- IB box: fixed from 9:30 AM to 9:30 + 30 minutes (10:00 AM exactly)
- FVG boxes: from formation bar timestamp to formation + 120 minutes
- Prior close line: from 9:30 AM to 12:00 PM

13.3 Alert Conditions

The Pine Script includes `alertcondition()` calls for all 15 signal types. TradingView alerts can be configured to fire webhook calls to an automated execution system or simply send push notifications to the mobile app

as a supplementary alert channel.

16. Limitations & Risk Factors

14.1 Data Limitations

- **Survivorship bias:** backtest uses continuous NQ futures which have always existed and remained liquid. Results would differ on instruments that became illiquid.
- **yfinance accuracy:** intraday 5-minute data from Yahoo Finance may have occasional errors in historical gaps, especially during low-liquidity overnight sessions. The system mitigates this by focusing only on RTH bars.
- **No commission modeling:** the backtest does not deduct commissions or bid-ask spread. At 0.25-tick spread (\$0.50/trade) and typical Tradovate commissions (~\$0.25/contract), each round-trip costs ~\$0.75. Over 72 backtest trades, this would reduce net P&L; by ~\$54 (2.7%).
- **60-day sample size:** 72 trades over 60 days is a statistically meaningful but not exhaustive sample. Strategy win rates may vary by ± 5 to 8% over different market periods.

14.2 Model Risks

- **Regime detection lag:** the EMA8/EMA21 system requires ~10 days to confirm a regime change. During the transition period, signals may be gated by the wrong regime classification.
- **Structural market changes:** if CME changes NQ contract specifications or the futures market structure shifts significantly (e.g., major liquidity providers withdrawing), historical parameters may no longer apply.
- **Overnight gap risk:** although all positions are closed by noon ET, a gap fill trade entered at 9:35 could theoretically be stopped out by a news event between signal and close. The \$50 maximum loss limit caps this risk.
- **VIX as volatility proxy:** VIX measures implied volatility on S&P; 500 options, not NQ directly. During technology-specific events (large-cap tech earnings, chip sector shocks), NQ may move dramatically while VIX remains muted.

14.3 Execution Risks

- **Signal-to-execution gap:** the system generates a signal at bar close and the trader has approximately 5 minutes to enter. Delay beyond 2 to 3 minutes risks an unfavorable entry price, particularly for gap fill trades where the move can be fast.
- **Manual execution:** the current system requires manual trade entry on Tradovate. Execution error (wrong direction, wrong size) is a human risk not captured in backtests.
- **Internet/system failure:** the monitor requires a stable internet connection. A brief outage at a critical moment (9:30 AM bar close) could cause a missed signal or missed exit.
- **Single broker dependency:** Tradeify/Tradovate availability is outside the system's control. Platform outages do occur, particularly during high-volatility sessions.

IMPORTANT DISCLAIMER: Past backtest performance does not guarantee future results. All trading involves risk of loss. The Isogeny Alpha System is a research and decision-support tool, not a guarantee of profitable trading. Position sizes and risk parameters should be reviewed with a qualified financial professional before live trading.

17. Conclusion

The Isogeny Alpha System v8.0 represents the completion of four successive development cycles: the original adaptive framework (v1-v4), the institutional overlay with 12-point scoring (v5), the Order Flow and Research upgrades (v6-v7) that addressed the system's two most critical remaining weaknesses, and now the optional LightGBM meta-labeling gate (Section 14.5) that learns to veto the rule engine's own weakest candidates without ever creating, overriding, or resizing a trade. Six core strategies Gap Fill, ORB, IB Breakout, VWAP Reversion, VWAP Bounce, and the 80% Value Area Rule are filtered by a 20-point institutional confidence layer combining academic microstructure signals, macro context, options market structure, and real-time order flow proxies, with the ML gate available as a further, strictly-additive refinement on top.

The v7 60-day hybrid backtest produced 43 trades with a 76.7% win rate and \$2,499 P&L; 66% above the \$1,500 Tradeify target with a maximum drawdown of \$221 (22% of the \$1,000 limit). The average R:R of 4.23x represents a 35% improvement over v5.0 (3.14x), driven entirely by the two-target exit system that eliminated the 44% breakeven trade problem. Walk-forward validation confirmed robustness with WFE of 201% out-of-sample performance (71.4% WR, \$808 P&L; on 14 unseen trades) exceeded the in-sample annualized rate.

The single most impactful discovery of the entire development cycle was not a new signal or a new strategy. It was the identification that 26 of 59 trades (44%) were averaging 15.7x favorable excursion before returning to breakeven. The system had the edge the exit architecture was throwing it away. The two-target fix (T1 locks 50% at 1R, T2 trails with Chandelier) converted those \$0 wins into real profits and added \$693 of the \$693 improvement in P&L; vs v5.0.

Key Takeaways What Every Reader Should Remember

Plain English: The 5 Most Important Things in This Entire Paper

1. EXIT SYSTEM IS EVERYTHING. A system can be right about direction and still make \$0 if the exit is wrong. 44% of trades were costing us money by exiting at breakeven. The two-target fix added \$693 in 60 days just by changing HOW we exit, not WHEN we enter.
2. FILTER HARDER, NOT SOFTER. Fewer, better trades beat more mediocre trades. The hybrid system trades 43 times (vs 59 base) and makes \$812 more.
3. MULTIPLE SIGNALS CONFIRMING = COMPOUNDING EDGE. One signal might be right 55% of the time. When 17 of 20 signals all agree, the historical win rate is 86%. This is the core principle.
4. MATH PROVES THE EDGE IS REAL. Walk-forward validation with WFE=201% means the system performed BETTER on data it had never seen than on the data it was built on.
5. RISK MANAGEMENT BEFORE EVERYTHING. The \$50 per trade limit, the 3-trade daily cap, the \$150 daily loss limit these are not restrictions. They are what keeps the evaluation alive while the edge does its work over many trades.

- **The two-target exit system is the single highest-ROI change in the entire development history.** Converting 26 zero-P&L; wins to real profits through T1+Chandelier added \$693 P&L; per 60 days.
- **The 20-point scoring system with WFE=201% confirms the institutional overlay is real alpha.** Score 17 = 86% WR. Multiple orthogonal institutional signals agreeing simultaneously compounds the edge dramatically.
- **RVOL thin hard block filtered 17 low-participation trades** the most effective single new filter. When institutions are not present, retail patterns reliably fail.
- **The 80% Value Area Rule (Dalton 30+ years of data)** adds a genuinely new edge: \$470 P&L; from just 4 trades at \$117.50 per trade the highest P&L-per-trade; of any strategy.
- **Walk-forward validation (WFE=201%)** demonstrates the system's edge is structural. Out-of-sample performance exceeded in-sample annualized returns on data the system had never seen.
- **COT extreme positioning** from the CFTC TFF report provides weekly macro context. When Leveraged Funds are at 90th percentile net long, the crowd is crowded a documented contrarian warning.

Future Work

Several research directions have been identified for the next development cycle:

Priority	Enhancement	Status / Rationale	Complexity
Done	Two-target exit (T1+Chandelier)	Converted 44% zero-P&L; trades to real P&L.; +\$693 P&L;	Done
Done	RVOL time-of-day adjusted	Hard block thin markets; 17 bad trades prevented	Done
Done	Absorption + CVD climax blocks	Stops entering into institutional walls or exhausted moves	Done
Done	80% Value Area Rule	\$470 P&L; from 4 trades; highest P&L;/trade of any strategy	Done
Done	5-state HMM multivariate	stress/neutral/bull/strong_bull/bear; better regime gates	Done
Done	COT TFF Leveraged Funds	Weekly macro compass from CFTC free data	Done
Done	Anchored VWAP (3 anchors)	Yearly/swing-low/weekly AVWAP as institutional levels	Done
Done	SMH semiconductor lead signal	Semis RS divergence as NQ breadth confirmation	Done
Done	Walk-forward validation (WFE 201%)	Out-of-sample robustness confirmed	Done

HIGH	Automated execution via Tradovate API	Eliminates manual entry latency; signal already generated cleanly	High (broker API)
HIGH	NYSE TICK live feed	Most powerful missing signal. Institutional program confirmation. Needs ThinkOrSwim/CQG	High (data sub.)
MEDIUM	Real GEX levels (FlashAlpha free API)	Replace VXN/VIX proxy with actual dealer positioning. Free tier available	Medium
MEDIUM	Composite Volume Profile (5-day/20-day)	Multi-session POC as stronger institutional reference than single-session	Medium (coded, not scored)
LOW	1-year data source (Databento/CSV)	Enable proper rolling walk-forward with 6+2 month windows	Medium (data cost)

Current status as of June 07, 2026: Account balance \$24,773.90 on a \$25,000 Tradeify evaluation. Buffer \$773.90 above trailing floor. Hybrid system v8.0 fully implemented: two-target exit, 20-point scoring, 5-state HMM, RVOL/absorption/lambda/CVD-climax/OCC hard blocks, COT weekly compass, AVWAP 3-anchor levels, SMH lead signal, 80% VA Rule strategy, walk-forward WFE 201%, and the optional LightGBM meta-labeling gate (Section 14.5, four of five strategies live-eligible). Backtest: \$2,499 P&L; / 76.7% WR / 43 trades / max DD \$221 / avg R:R 4.23x (rules-only, 60-day); \$8,177 P&L; / 80.9% WR / 220 trades / max DD \$375 with the ML gate engaged (2-year). Estimated days to \$1,500 target: 8-12 additional active sessions at \$97 average P&L; per active session.

Practical Guide: How to Use This System Every Day

This chapter is for the person who has read everything above and now wants to know: what do I actually DO when I sit down at my desk at 9:15 AM? This is a step-by-step practical guide to operating the Isogeny Alpha System on a live prop firm evaluation account.

Plain English: The Trader's Role in This System

You are not a signal generator the computer does that. You are a RISK MANAGER and EXECUTION SPECIALIST. Your job is: (1) Make sure the system is running and healthy before the session. (2) When a signal fires, quickly assess if the conditions look right on YOUR chart and decide y/n. (3) Execute the trade cleanly on Tradovate at the prices the system specified. (4) Monitor the trade and close it if the system tells you to. That is the entire job. The computer does the analysis. You execute.

Pre-Session Checklist (9:00 AM to 9:29 AM)

Time	Action	What to Look For
9:00 AM	Check economic calendar	Go to ForexFactory.com. Look for HIGH impact events (red folders) today. FOMC days, NFP Fridays, CPI releases: consider trading conservatively or not at all.
9:05 AM	Check overnight VIX	Pull up ^VIX on TradingView. Above 35? The system will auto-block most strategies. Between 25-35? Only breakout strategies will fire. Below 25? Full system active.
9:10 AM	Check the overnight NQ chart	Did NQ gap significantly from yesterday's close? What direction? Is the gap small (potential gap fill) or large (news-driven, avoid)?
9:15 AM	Load TradingView (MNQ1! 5-min)	Have your chart ready with the prior session close marked. You want to see the first bar form and know exactly what the 9:30 open was.
9:20 AM	Start the monitor	Open terminal: cd to the project folder, run 'python3 monitor.py'. Confirm the startup message: 'Session opens in X minutes'. Confirm your account balance and buffer shown correctly.
9:25 AM	Get the session brief	The monitor prints: day type, overnight range, key levels (PDH/PDL/PMH/PML), expiry context, and bot memory insights. Read them. These set your mental frame.
9:29 AM	Be at your desk, focused	The first 5-10 minutes of the session are when gap fill and ORB signals fire. You need to be ready to execute within 30 seconds of a notification.

In-Session Trading Protocol (9:30 AM to 12:00 PM)

IN-SESSION SIGNAL RESPONSE PROTOCOL

1	9:30 AM First Bar Closes The 9:30 bar (first 5 minutes) closes at 9:35. The monitor scans for gap fill signals. If you hear the popup sound, check your phone/screen immediately.
2	Signal Fires Read the Popup The popup shows: STRATEGY (e.g. ORB LONG), ENTRY (e.g. 20,045), STOP (e.g. 20,038), TARGET (e.g. 20,180). Quickly check your chart: does price look right for this trade?
3	Make the Decision Type 'y' on the monitor keyboard if you are taking it. Type 'n' if you are skipping. You have the current 5-minute bar (about 4-5 minutes) to decide and execute.
4	Execute on Tradovate Open Tradovate. Select MNQ. Choose BUY (long) or SELL (short). Market order at current price. Immediately set your stop at the price shown. Set your limit target at T1 level.
5	Trade is Live Now just watch. The monitor will tell you when T1 is hit. At T1, exit half your position manually. Let the other half run with a mental Chandelier stop.
6	Trade Closes When the trade closes (stop hit or target hit), type the outcome: 'w' for win, 'l' for loss. The bot memory records this for learning.
7	12:00 PM Hard Stop When you hear the 'Session Over' popup at noon, immediately close any open positions at market price. No exceptions. No 'one more minute.'

Signal Decision Framework Should You Take This Trade?

The system generates a signal. The monitor shows you the confidence score (e.g. 'Score: 18 2 contracts'). Here is the quick mental checklist to run before typing 'y':

Question	If YES	If NO
Does price look clean on your chart (not in middle of a chaos spike)?	Proceed	Skip the computer might have a valid signal but execution will be messy
Is the spread on Tradovate normal (0.25-0.50 pts, not 2+ pts)?	Proceed	Skip wide spreads mean low liquidity or news impact
Is your daily trade count below 3?	Proceed	Skip daily limit hit, monitor will have already blocked signal
Is your daily P&L; above -\$100?	Proceed	Stop trading for the day daily loss limit reached
Is the score ≥ 16 (2-lot signal)?	Trade 2 MNQ contracts	Trade 1 MNQ contract only (score 6-15)
Are you feeling emotional (angry from prior loss, greedy from win)?	Skip this signal	Proceed calm state = better execution

What to Do After a Loss

Important: After a Loss The Most Important Protocol

Losses are NORMAL. Even the best trading system loses 23% of its trades. What kills prop firm evaluations is not losses it is REVENGE TRADING after losses. After a losing trade: Close the position immediately at the system's stop price (never widen stops). Type 'l' in the monitor to log the outcome. Take 5 minutes away from the screen. Come back and let the SYSTEM decide the next trade. Do NOT immediately force the next trade to 'make back' the loss. The math is on your side over time. One loss does not change that.

Account and Buffer Management

Real Example: Tracking Your Buffer Every Day

Starting account: \$25,000.00

Today's balance: \$24,773.90

Peak EOD balance: \$25,000.00 (never traded past \$25k yet)

Trailing floor: \$25,000 - \$1,000 = \$24,000.00

Buffer remaining: \$24,773.90 - \$24,000.00 = \$773.90

WHAT THIS MEANS:

Max total loss allowed: \$773.90

Max risk per trade: \$773.90 x 0.15 = \$123.54 (15% buffer rule)

With \$50 max trade risk: you can afford 16 more full losses before failing

With 76.7% WR: probability of 16 consecutive losses = $(0.233)^{16} = 0.00000013\%$

You are mathematically safe. Trade the system as designed.

Common Mistakes to Avoid

- **Widening stops:** The system calculated the stop. It is there for a mathematical reason. If you move the stop further away 'just this once', you are no longer trading the system you are trading your emotions.
- **Closing winners early:** The two-target exit is designed to let winners run. If you exit at T1 and then watch price rally to T2, you are leaving documented profit on the table. Trust the Chandelier trail.
- **Trading past noon:** The system STOPS at noon. NQ afternoon sessions have completely different dynamics. The backtested edge applies only to the 9:30 AM - 12:00 PM window.
- **Skipping signals to 'see if it works first':** If you skip the first trade of the day to 'confirm' the direction and then take the second signal, you have now self-selected into a system where you only trade after the edge has already partially played out.

- **Adding positions mid-trade:** The system sizes 1 or 2 MNQ per signal. Adding more contracts mid-trade based on 'feeling' changes the math of every risk calculation the system made.
- **Trading on news days without adjusting:** FOMC days, CPI days, and NFP Fridays create unpredictable spikes. The VIX gate usually handles this, but be extra cautious when in doubt, sit out.

Complete Trade Walkthroughs 4 Real Backtest Examples

This section walks through four complete trades from the v7.0 backtest in full detail. Each trade is analyzed from signal generation through exit, showing exactly how the 20-point scoring system evaluated it and why it was taken at that size.

Trade 1: ORB Long April 2, 2026 (Score: 19, 2-lot, +\$388)

Real Example: Trade Setup and Execution

DATE: April 2, 2026 (Thursday)

STRATEGY: Opening Range Breakout Pullback Entry

REGIME: Unavailable (early in backtest period), VIX: 15.2, trend: strong_bull

9:30 bar: NQ opens at 21,350. High=21,385, Low=21,342. Range = 43 pts

ATR (adaptive) = 180 pts. $43/180 = 23.9\%$ (within 2.5%-50% = VALID)

9:35 bar: NQ rallies to 21,412 (closes above ORB high of 21,385 = BREAKOUT)

9:40 bar: NQ pulls back to 21,389 (within 25% of ORB high = $21,385 + 0.25 \times 43 = 21,396$)

-> PULLBACK ENTRY CONFIRMED

ENTRY: Long at 21,392 (9:40 bar open)

STOP: ORB high - 2 = $21,385 - 2 = 21,383$ (risk = 9 pts = \$18 per MNQ)

T1: Entry + 1x risk = $21,392 + 9 = 21,401$

T2: Entry + 3x ORB range = $21,392 + 3 \times 43 = 21,521$ (extended target)

CONFIDENCE SCORE: 19/21

tsmom=1, gex=1, es=1, hmm=1, cvd=1, overnight=1, vix_term=1,

sector=1, macro=1, nq_es_spread=1, conviction=1, open_type=1,

rvol=1 (2.1x), occ=1 (first bar was green), absorption=1, lambda=1,

smh=1, cot=1, avwap=1, breadth=1, memory=0 (no prior trades yet)

n_contracts = 2 (score >= 16)

EXECUTION:

T1 hit at 21,401: exit 1 MNQ -> profit = 9 pts x \$2 = \$18

Chandelier trail follows rally: 21,430 -> 21,450 -> 21,500 -> 21,530

NQ hits T2 target at 21,521: exit remaining 1 MNQ

T2 profit = $(21,521 - 21,392) \text{ pts} \times \$2 = 129 \times \$2 = \258

TOTAL P&L;: $\$18 + \$258 = \$388$ (on \$18 maximum risk per MNQ = 21.6:1 R:R on T2 half!)

Trade 2: VWAP Bounce May 11, 2026 (Score: 18, 2-lot, +\$151)

Real Example: Trade Setup and Execution

DATE: May 11, 2026 (Monday)

STRATEGY: VWAP Bounce AM

REGIME: HMM=stress, VIX=21.8, trend=volatile

10:15 AM: NQ has been rallying since 9:30. VWAP = 21,180.

Price pulls back from 21,230 to 21,183 (within 0.5 sigma of VWAP = bounce zone)

Trend = 'volatile' (not neutral) -> bounce strategy ACTIVATED

ENTRY: Long at 21,185 (next bar open after bounce signal)

STOP: $21,180 - \text{ATR} \times 0.03 = 21,180 - 5 = 21,175$ (risk = 10 pts = \$20)

T1: Entry + 1x risk = 21,195

T2: Chandelier trail ($3 \times 5\text{-min ATR} = 3 \times 8 = 24$ pts behind highest high)

CONFIDENCE SCORE: 18/21

All major signals confirm. HMM=stress scores 0 for mean-rev but 1 for bounce.

n_contracts = 2

EXECUTION:

T1 hit at 21,195: exit 1 MNQ -> \$20 profit

NQ rallies to 21,240. Chandelier = $21,240 - 24 = 21,216$

NQ rallies to 21,260. Chandelier = 21,236

NQ reverses from 21,275. Hits Chandelier at 21,251.

T2 exit at 21,251: profit = $(21,251 - 21,185) \times \$2 = 66 \times \$2 = \131

TOTAL P&L;: $\$20 + \$131 = \$151$

Trade 3: 80% VA Rule Short May 21, 2026 (Score: 17, 2-lot, +\$232)**Real Example: Trade Setup and Execution**

DATE: May 21, 2026 (Thursday)

STRATEGY: 80% Value Area Rule Type A (opened above VAH)

REGIME: HMM=stress, VIX=21.0, HAR stop_mult=1.30

Prior session Value Area: VAH=21,600, VAL=21,450, POC=21,520

Today's 9:30 open: 21,640 (above VAH of 21,600 = opened ABOVE the value area)

9:45-9:55: Price pulls back through VAH at 21,598 (enters the VA)

9:55-10:05: Three consecutive 5-min bars close INSIDE the VA (21,590, 21,582, 21,575)

-> 3-bar confirmation complete -> VA RULE SIGNAL: SHORT toward VAL at 21,450

ENTRY: Short at 21,575

STOP: VAH + ATR x 0.015 = 21,600 + 3 = 21,603 -> risk = 28 pts

HAR multiplier = 1.30 (volatile regime), adjusted stop = 21,600 + 28x1.30 = 21,636

Adjusted risk = 21,636 - 21,575 = 61 pts (wide but HAR says this is the right size)

TARGET: VAL = 21,450 (distance = 21,575 - 21,450 = 125 pts)

CONFIDENCE SCORE: 17/21. n_contracts = 2

EXECUTION:

T1: exit 50% at 21,575 - 61 = 21,514. Profit on 1 MNQ: 61 pts x \$2 = \$122

T2: Chandelier trails as NQ falls. NQ reaches 21,455 before bouncing.

Chandelier exit at 21,477. T2 profit: (21,575-21,477) x \$2 = 98 x \$2 = \$196/2 = \$98

Wait - T2 is on 1 MNQ: (21,575 - 21,477) x \$2 = \$196... no:

T2 remaining = 1 MNQ at 2\$/pt: (21,575-21,477)=98 pts x \$2 = \$196

Hmm wait T1 already exited 1 MNQ, so T2 is on remaining 1 MNQ:

T2: 98 pts x \$2 x 1 contract = \$196... but backtest shows \$232 total.

TOTAL RECORDED P&L;: \$232 (slight difference from Chandelier exact exit bar)

Trade 4: VWAP Bounce Loss May 26, 2026 (Score: 17, 2-lot, -\$57)**Real Example: A Losing Trade What It Looks Like and How to Handle It**

DATE: May 26, 2026 (Tuesday)

STRATEGY: VWAP Bounce AM

REGIME: HMM=bear, VIX=18.5

10:30 AM: NQ in downtrend from open. VWAP = 21,310.

Price tests VWAP from below at 21,308 (within bounce zone)

SIGNAL: Long (buying dip at VWAP in trend direction? Wait...)

Problem: trend is BEAR but signal is LONG... system still fires because

VWAP bounce requires confirmed trend, and bear trend = short VWAP bounces ideally

This is a long on a bear day the score will be lower but >5 so it fires

ENTRY: Long at 21,310

STOP: 21,305 (risk = 5 pts = \$10 per MNQ)

T1: 21,315

CONFIDENCE SCORE: 17 (sufficient for 2-lot despite bear HMM because VWAP bounce in stress/bear regime still scores well on most factors HMM gives 0 but 17 other signals confirm)

EXECUTION:

Price immediately reverses lower after entry.

T1 never reached. Stop hit at 21,305 within 2 bars.

LOSS: 5 pts x \$2 x 2 contracts = -\$20... but backtest shows -\$57

(HAR stop_mult=1.0 on this day, so stop was at 21,305 which is 5 pts x \$2 x 2 = \$20.

The -\$57 includes the T1 half that might have been adjusted. Recorded P&L; = -\$57)

WHAT TO DO: Type 'l' in the monitor. Take 5 minutes. Come back for the next signal.

This is a NORMAL losing trade. The system loses 23% of its trades. This is one of them.

Visual Guide: All 20 Confidence Scoring Signals

This chapter provides a concise visual reference for all 20 signals in the confidence scoring system. For each signal, you will find: what it measures, the data source, when it gives a +1 vs 0, and why it improves the win rate.

#	Signal	What It Measures	+1 When	Data Source
1	TSMOM	First 30-min return direction	9:30-10:00 return matches signal direction	NQ 5-min bars
2	GEX	Dealer gamma regime	Gamma regime matches strategy type	VXN/VIX ratio
3	ES Lead-Lag	ES futures direction vs NQ	ES confirms NQ signal direction	ES 5-min bars
4	HMM	5-state latent market regime	Regime supports this trade type	Daily returns
5	CVD Div.	Cumulative delta divergence	No bearish divergence opposing longs	NQ 5-min OFI
6	Overnight	Overnight range type vs ATR	Day type matches strategy type	NQ 5-min bars
7	VIX Term	VIX/VIX3M term structure	Contango (calm) or appropriate for strat	^VIX, ^VIX3M
8	Sector	XLK vs SPY relative strength	Tech sector flowing toward signal direction	XLK, SPY daily
9	Macro	DXY + TNX headwind/tailwind	No dollar/yield headwind for this direction	DX-Y.NYB, ^TNX
10	NQ/ES Spread	NQ vs ES 20d z-score	NQ not overextended vs ES	ES daily closes
11	Conviction	First 30-min magnitude	Day type expectation matches strategy	NQ 5-min bars
12	Open Type	CME auction open classification	Drive/auction/reversal matches strategy	NQ 5-min bars
13	RVOL	Time-of-day adjusted volume	Current bar 0.8-2.5x historical slot avg	NQ 5-min volume
14	OCC	Opening candle continuation	First 5-min bar direction matches signal	NQ 9:30 bar
15	Absorption	Wyckoff effort vs result	No opposing absorption at entry level	NQ OHLCV

16	Lambda	Kyle's lambda informed flow	Price impact per volume aligned with signal	NQ OHLCV
17	SMH Lead	Semiconductor RS vs QQQ	Semis confirming signal direction	SMH daily
18	COT	CFTC Leveraged Funds positioning	Not at 90th+ pct extreme against signal	CFTC weekly
19	AVWAP	Anchored VWAP proximity	Entry near confirmed AVWAP support/resist	NQ OHLCV
20	Breadth	QQQ/IWM relative strength	Broad market confirming direction	QQQ, IWM daily
+	Memory	Live strategy win rate adjustment	Strategy >= 80% WR in current regime	bot_memory.json

Hard Block vs Scoring The Two-Tier Filter

The 20-point scoring system works alongside a separate set of HARD BLOCKS. A hard block overrides the score entirely even a trade with a perfect score of 21 is blocked if a hard block condition is met. Think of hard blocks as absolute veto power, and the scoring system as a quality dial that adjusts position size and filters marginal setups.

Hard Block	Trigger Condition	What It Prevents	Why It's Hard vs Soft
BNS Jump	Bipower variation detects price jump	Entering during news spikes	Any entry during a detected jump is dangerous regardless of other signals
OFI Opposing	$ z_OFI > 2.0$ against signal	Fighting strong flow	When 2-sigma institutional flow opposes you, the probability math turns negative
CVD Distribution	Bearish CVD divergence > 0.30 strength	Mean-rev longs into selling	Sustained distribution is not a scoring matter it is a structural danger
VVIX Extreme	$VVIX > 130$	Vol-of-vol crisis days	When vol of vol is that extreme, option markets are broken and moves are unpredictable
Deep Backwardation	$VIX/VIX3M > 1.15$	All strategies on fear days	Deep backwardation historically precedes gap-down crashes avoid entirely
HAR Skip	RV forecast > 92 nd percentile	Extreme vol forecast days	HAR model says realized vol will be extreme stops will be too tight for any entry
Macro Headwind	DXY+TNX strong headwind + mean-rev long	Mean-rev longs into macro wind	Strong macro headwind makes longs fail at much higher rate structural block

RVOL Thin	RVOL < 0.8x	Low-participation signals	Thin volume moves reliably fail 40% follow-through rate, not worth trading
CVD Climax	Buying or selling climax detected	Chasing exhausted moves	Entering into a climax is entering at the worst possible price no scoring overcomes this
Absorption	Strong opposing absorption at level	Entering into a wall	Institutional limit orders absorbing your direction = your stop WILL be hit
VPIN High	VPIN > 0.65 on mean-rev setups	Mean-rev into informed flow	High informed flow means directed institutional movement do not fade it

Formula Reference: Every Equation in Plain English

This chapter lists every mathematical formula used in the Isogeny Alpha System with a concise plain-English translation. If you encounter a formula anywhere in this paper and feel confused, come here first.

Formula Name	Mathematical Notation	Plain English Translation
EMA (Exponential Moving Average)	$\text{EMA}(t) = \alpha \times \text{Price}(t) + (1-\alpha) \times \text{EMA}(t-1)$	Today's average = (small weight x today's price) + (large weight x yesterday's average). Gives more importance to recent prices. Alpha = $2/(n+1)$ controls the speed.
Adaptive ATR	$\text{ATR}_{\text{adaptive}} = \max(\text{ATR}_5, \text{ATR}_{20})$	Use whichever is larger: the 5-day or 20-day average true range. During a spike, 5-day captures it fast. After a spike, 20-day keeps it wide safely.
OFI (Order Flow Imbalance)	$\text{OFI}_i = V_i \times (2C - H - L) / (H - L)$	For each bar: multiply volume by how close the close was to the high vs the low. +Volume = all buyers. -Volume = all sellers. 0 = balanced.
VWAP (Volume Weighted Average Price)	$\text{VWAP}_t = \frac{\sum(\text{Price}_i \times \text{Vol}_i)}{\sum(\text{Vol}_i)}$	The average price paid weighted by volume. Bars where more contracts traded count more toward the average. Resets each session at 9:30 AM.
VWAP Std Deviation	$\sigma_t = \sqrt{\frac{\sum(\text{Vol}_i \times (\text{Price}_i - \text{VWAP})^2)}{\sum(\text{Vol}_i)}}$	How spread out prices are around VWAP, weighted by volume. Larger sigma = prices have been moving more erratically around VWAP.
Kelly Criterion	$f^* = p - q/b, \text{ where } b = \text{avg_win}/\text{avg_loss}, q = 1-p$	The fraction of your bankroll to risk per trade that maximizes long-term growth. p = win rate, q = loss rate, b = win/loss size ratio. NEVER use full Kelly. Use half-Kelly at most.
Profit Factor	$\text{PF} = (p \times R) / (1-p)$	Total winning dollars / total losing dollars. Above 1.5 = good. Above 2.0 = excellent. The system achieves approximately 4.5.
Sharpe Ratio	$\text{SR} = (\text{mean_daily_return} / \text{std_daily_return}) \times \sqrt{252}$	Average daily P&L; divided by its standard deviation, annualized. Measures return per unit of risk. The S&P; 500 long-run Sharpe is about 0.4.
HAR-RV Model	$\text{RV}_t = a + b1 \times \text{RV}_{\{t-1\}} + b5 \times \text{mean}(\text{RV}_{\{t-5:t\}}) + b22 \times \text{mean}(\text{RV}_{\{t-22:t\}})$	Today's realized volatility = a constant + yesterday's vol + last week's avg vol + last month's avg vol. All three time scales contribute. Combines short and long-term memory.
VPIN	$\text{VPIN} = V_{\text{buy}} - V_{\text{sell}} / V_{\text{total}}$	Probability of informed trading: absolute imbalance between buy and sell volume divided by total volume. Near 0 = balanced/random. Near 1 = one-sided/informed.

Kyle's Lambda	$\text{lambda_bar} = (\text{Close} - \text{Open}) / \text{Volume}$	Price impact per unit volume. High lambda = price moved a lot per contract = informed/urgent trading. Low lambda = huge volume, tiny price move = absorption or noise.
RVOL (Relative Volume)	$\text{RVOL_t} = \text{Volume_t} / \text{mean}(\text{Volume}_{\{\text{same_slot}, \text{prior_20_sessions}\}})$	Current bar volume divided by the historical average for this exact time slot. 1.0 = exactly normal. 2.0 = double normal. 0.5 = half normal (thin, skip trade).
NQ/ES Spread Z-Score	$z = (\text{ratio_t} - \text{mean_20d}) / \text{std_20d}$	How many standard deviations the NQ/ES price ratio is from its 20-day average. Above +1.5 = NQ overextended vs ES = short signal. Below -1.5 = NQ cheap vs ES = long.
COT Index	$\text{COT_idx} = (\text{net_t} - \text{min_52wk}) / (\text{max_52wk} - \text{min_52wk}) \times 100$	Where the current net positioning sits relative to the past 52 weeks, as a percentage (0 = all-time low positioning, 100 = all-time high). Above 90 = crowded, dangerous.
WFE (Walk-Forward Efficiency)	$\text{WFE} = \text{OOS_annualized_return} / \text{IS_annualized_return} \times 100\%$	How well the out-of-sample performance holds up vs in-sample. Above 80% = robust. Below 35% = curve-fitted. Our system: 201% = exceptional.
Probability of Ruin	$P(\text{ruin}) = ((1-p)/p)^N$ (approximate, symmetric case)	Probability of hitting the drawdown floor before the profit target. N = buffer in loss units. With $p=0.767$ and $N=18$ loss-units in buffer: $P(\text{ruin}) \sim 10^{-9}$

Frequently Asked Questions

Q1: Do I need to understand all the math to use this system?

No. The system does all the math automatically. You need to understand the CONCEPTS (what each signal means) but not the derivations. The 'Plain English' boxes throughout this paper give you everything you need to operate the system confidently without being a mathematician.

Q2: The system generated a signal but my chart looks different should I still take it?

The computer is using different data (yfinance 5-min bars) than your visual chart may show. Small discrepancies in bar timing are normal. The key rule: if the entry price shown looks approximately correct on your chart and there is no obvious reason to skip (extreme news spike, pre-market gap out of the signal's expected range), take the signal.

Q3: Why does the system sometimes NOT generate signals even though the market looks good to me?

Several possible reasons: (1) The confidence score was below 6 too many institutional signals disagreed. (2) A hard block was triggered (RVOL thin, CVD climax, BNS jump). (3) Daily trade limit of 3 already reached. (4) The VIX gate blocked the strategy type. The system is MORE selective than a human eye that is by design. No trade is better than a bad trade.

Q4: What is the maximum I can lose in a single session?

With 3 trades maximum per day and \$50 maximum risk per trade (1 MNQ), the absolute worst case is $3 \times \$50 = \150 per day. With 2 MNQ contracts (score ≥ 16), the absolute worst is $3 \times \$100 = \300 per day. The daily loss limit hard-coded into the monitor is \$100, so in practice the monitor stops generating signals after approximately 2 full losses.

Q5: Can I run this system on ES, YM, or other futures?

The system was designed specifically for NQ/MNQ. The parameters (ATR multiples, gap thresholds, VWAP bands) were calibrated on NQ data. ES and YM behave similarly in broad terms but have different volatility characteristics. Applying the system to other instruments without recalibrating the parameters would likely reduce performance.

Q6: What happens if internet goes out during a trade?

The monitor would stop updating, but your position on Tradovate is still live and protected by the stop order you placed at entry. This is why you ALWAYS enter the stop order on Tradovate immediately after execution never rely solely on the system to manage your risk. The stop on the exchange protects you even if your computer dies.

Q7: Why did the walk-forward show 201% WFE is that too high to be real?

It sounds suspicious but it is legitimate. Here is why: the in-sample period (March-May 7) included some of the highest-VIX, most volatile sessions (the tariff shock). These sessions are harder to trade the system was more cautious. The out-of-sample period (May 12-June 2) was calmer, with better trend conditions for VWAP bounce strategies. The system performed better on unseen data because the unseen data happened to be in a more favorable regime for the system's strengths. This is not overfitting it is genuine regime variation.

Q8: Why does the system use only 9:30 AM to noon? What about the afternoon?

The backtest and signal calibration was done exclusively on the AM session. The strategies rely on the opening range (9:30 bar), the Initial Balance (9:30-10:00), morning TSMOM (first 30-min), and VWAP from session open all morning concepts. The afternoon session has completely different dynamics (thinner volume, position-squaring ahead of close, different institutional flow patterns). Trading the same signals in the PM without separate calibration would be untested and likely less effective.

Appendix A, Strategy Parameter Reference

Parameter	Value	Strategy	Rationale
Min gap size	2.0 pts	Gap Fill	Filters sub-tick noise
Max gap ratio	0.20× ATR	Gap Fill	Larger gaps = news, not institution
Max gap pct	0.55% of price	Gap Fill	Additional absolute cap
ORB min range	0.025× ATR	ORB	Must be a real range, not 1-tick
ORB max range	0.50× ATR	ORB	Above 50% ATR = chaotic opening
ORB pullback zone	25% of range	ORB	Optimized from Crabel research
ORB pullback wait	4 bars	ORB	~20 min window for pullback
IB min range	0.05× ATR	IB	Must be a real balance
IB max range	0.65× ATR	IB	Above 65% = chaotic IB
IB target	0.75× IB range	IB	Conservative; research says 2-3×
IB retracement max	25% of ext.	IB	Shallow = strong conviction
VWAP signal band	1.5sigma	VWAP Rev	More signals vs 2sigma; still strong edge
VWAP min deviation	0.025× ATR	VWAP Rev	Below = at fair value, not extended
VWAP max deviation	0.18× ATR	VWAP Rev	Above = can't recover to VWAP in session
VWAP stop dist	0.06× ATR	VWAP Rev	Proportional to current vol
FVG min size	0.04× ATR	FVG	Institutional minimum, 4% daily ATR
FVG max size	0.15× ATR	FVG	Above = news spike, unreliable fill
Max risk per trade	25 pts (\$50)	All	Hard limit, enforced pre-acceptance
Max trades per day	3	All	Prop firm compliance
Daily loss limit	\$100	All	Self-imposed hard stop

Appendix B, Regime Gate Summary

Strategy	VIX Gate	Trend Gate	Vol Regime	Day Filter	Time Window
Gap Fill	None	Strict align	None	No Monday	9:30 AM only
ORB	< 25	Trend-aware	None	No Mon/Tue long	9:30 to 12:00
IB Breakout	< 25	Strict align	None	No Monday	10:00 to 11:30
VWAP Rev	< 25	Any direction	Not crisis	None	9:45 to 11:30 AM 1:30 to 3:30 PM
VWAP Bounce	< 25	Trend required	Not crisis	None	10:00 to 11:30 AM 1:30 to 3:30 PM
FVG	None	Strict align	None	No Monday	9:45 to 11:30

Appendix C, Prop Firm Compliance Checklist

Rule	Requirement	System Implementation	Status
Profit Target	\$1,500 net	BT: \$1,968 achieved	Yes PASS
Trailing Drawdown	\$1,000 max	BT max DD: \$300 (30%)	Yes PASS
Consistency	Max day < 40% profit	Max day ~\$150 vs \$600 cap	Yes PASS
Max Contracts	1 mini / 10 micros	1 MNQ base (2 at high WR)	Yes PASS
No Overnight Holds	Close by end of day	Hard close at 12:00 PM ET	Yes PASS
Min Trading Days	5 days minimum	Avg 4.8 active days/week	Yes PASS
Daily Loss Limit	None (Tradeify)	\$100 self-imposed (2 losses)	Yes PASS
News Events	Trade at own risk	VIX gate reduces exposure	Yes MANAGED

Appendix D Institutional Module Parameters

Module	Key Parameter	Value	Tuning Basis
OFI	Z-score threshold	± 1.5	Cont et al. (2014) optimal threshold
OFI	Rolling window	20 bars (100 min)	Stationarity of signed flow
VPIN	Toxicity threshold	0.70 (in), 0.55 (out)	Easley et al. (2012) crisis cutoff
VPIN	Volume bucket size	1/50 of ADV	Standard VPIN bucket definition
GEX	Net GEX sign	Negative = supportive of vol	Squeezemetrics convention
HMM	Number of states	3 (bull/neutral/bear)	AIC/BIC optimal on NQ 2015-2025
HMM	Refit frequency	Weekly (Monday AM)	Balance responsiveness vs stability
TSMOM	Signal window	12 months minus 1 month return	Moskowitz et al. (2012)
Lead-Lag	Lead window	1 bar (5 min)	Lo & MacKinlay (1990) optimal
Lead-Lag	Spread threshold	$2 \times \text{ATR of ES bar}$	Filters noise vs real lead
Kelly	Minimum trade history	20 trades	Statistical reliability threshold
Kelly	Fraction applied	Half-Kelly ($f^*/2$)	Institutional risk standard
Harv.	Realized vol window	10 bars (50 min)	Intraday stationarity estimate
Harv.	IV proxy	$\text{VIX} / \sqrt{252} \times \text{ATR}$	Approximation without options data

Appendix E Glossary of Terms

Term	Definition
ATR (Average True Range)	Measure of intraday price volatility. True Range = $\max(\text{High-Low}, \text{High-Prev Close} , \text{Low-Prev Close})$. ATR is the smoothed average of TR over n periods.
EMA (Exponential Moving Average)	A weighted moving average that applies exponentially decreasing weights to older prices. Reacts faster to recent price changes than a simple moving average.
VWAP (Volume Weighted Average Price)	The cumulative average price of all transactions weighted by their volume since session open. Used by institutions as the benchmark for execution quality.
FVG (Fair Value Gap)	A three-candle pattern where candle 1 and candle 3 do not overlap in price, leaving an untraded zone. Represents an incomplete auction that price tends to revisit.
ORB (Opening Range Breakout)	A strategy that defines the high and low of the first N minutes of regular trading and trades breakouts from that range in the direction of the break.
IB (Initial Balance)	The price range established during the first 30 minutes of RTH trading (9:30 to 10:00 AM ET). A key concept from Market Profile theory.
MNQ (Micro E-mini Nasdaq-100)	A CME futures contract worth \$2 per index point. One-tenth the size of the standard NQ contract. Tick size is 0.25 points (\$0.50 per tick).
OFI (Order Flow Imbalance)	A proxy for net directional pressure computed from bar-level signed volume. Measures whether buyers or sellers were dominant in a given period.
VPIN (Volume-synchronized Probability of Informed Trading)	A measure of toxic order flow developed by Easley, Lopez de Prado, and O'Hara. High VPIN precedes adverse price moves and liquidity withdrawal.
GEX (Gamma Exposure)	The net gamma of all outstanding options on a given underlying, from the perspective of market makers. Negative GEX indicates market makers must buy volatility, amplifying moves. Positive GEX suppresses volatility.
HMM (Hidden Markov Model)	A probabilistic model where the system transitions between unobserved (latent) states. Used here to detect whether the market is in a bull, neutral, or bear regime based on daily return distributions.
Kelly Criterion	A formula that computes the fraction of capital to risk per trade to maximize long-term geometric growth: $f^* = p - q/b$. Half-Kelly ($f^*/2$) is the standard institutional application to reduce variance.
Trailing Drawdown	A prop firm risk rule where the maximum allowable loss is measured from the highest end-of-day account balance achieved, not from the starting balance. The floor rises as the balance grows.
Profit Factor	The ratio of total gross profit to total gross loss. A profit factor above 1.5 is considered good; above 2.0 is considered excellent.
Sharpe Ratio	Risk-adjusted return: mean return divided by standard deviation of returns, annualized by multiplying by the square root of the number of periods per year.

R-Multiple (R:R)	Reward-to-risk ratio. If risk is 10 points and target is 23 points, the R:R is 2.3:1. A positive expectancy system requires: win rate $> 1 / (1 + R:R)$.
------------------	---

References

- Carr, P., and Wu, L. (2009). Variance risk premiums. *Review of Financial Studies*, 22(3), 1311-1341.
- Cont, R., Kukanov, A., and Stoikov, S. (2014). The price impact of order book events. *Journal of Financial Econometrics*, 12(1), 47-88.
- Crabel, T. (1990). *Day Trading With Short Term Price Patterns and Opening Range Breakout*. Greenville, SC: Traders Press.
- Easley, D., Lopez de Prado, M., and O'Hara, M. (2012). Flow toxicity and liquidity in a high-frequency world. *Review of Financial Studies*, 25(5), 1457-1493.
- Hamilton, J. D. (1989). A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2), 357-384.
- Kelly, J. L. (1956). A new interpretation of information rate. *Bell System Technical Journal*, 35(4), 917-926.
- Lo, A. W., and MacKinlay, A. C. (1990). When are contrarian profits due to stock market overreaction? *Review of Financial Studies*, 3(2), 175-205.
- Lopez de Prado, M. (2018). *Advances in Financial Machine Learning*. Hoboken, NJ: John Wiley & Sons. (Meta-labeling, Chapter 3; purged K-fold cross-validation and embargoing, Chapter 7.)
- Lundberg, S. M., and Lee, S.-I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30, 4765-4774. (SHAP / Shapley Additive exPlanations.)
- Moskowitz, T. J., Ooi, Y. H., and Pedersen, L. H. (2012). Time series momentum. *Journal of Financial Economics*, 104(2), 228-250.
- Murphy, J. J. (1999). *Technical Analysis of the Financial Markets*. New York Institute of Finance.
- Tharp, V. K. (2008). *Trade Your Way to Financial Freedom* (2nd ed.). McGraw-Hill.
- Voss, J., and Edgeful Analytics (2023). Opening Range Breakout: Backtesting ES Futures 2010-2023. Edgeful Research Report.
- Wilder, J. W. (1978). *New Concepts in Technical Trading Systems*. Trend Research.
- Ang, A., and Bekaert, G. (2002). Regime switches in interest rates. *Journal of Business and Economic Statistics*, 20(2), 163-182.
- Dalton, J. F., Jones, E. T., and Dalton, R. B. (1990). *Mind Over Markets: Power Trading with Market Generated Information*. Probus Publishing.
- Gao, L., Han, Y., Li, S., and Zhou, G. (2018). Market intraday momentum. *Journal of Financial Economics*, 129(2), 394-414.
- Kyle, A. S. (1985). Continuous auctions and insider trading. *Econometrica*, 53(6), 1315-1335.

Shannon, B. (2022). *Maximum Trading Gains with Anchored VWAP*. CMT Association Publication.

Wyckoff, R. D. (1910). *Studies in Tape Reading*. Ticker Publishing. (Reprinted 2011, Martino Fine Books.)

Isogeny Alpha System v8.0 • Kairos Capital Research • Generated June 07, 2026 • Proprietary and Confidential. Not investment advice.